

HIDE and Seek: Detecting Hallucinations in Language Models via Decoupled Representations

Anwoy Chatterjee*, Yash Goel, Tanmoy Chakraborty

Indian Institute of Technology Delhi, India

*Contemporary language models can generate fluent responses that are factually incorrect or unfaithful to the input context. We propose HIDE, a training-free, single-pass hallucination detector that applies an adapted Hilbert–Schmidt Independence Criterion (HSIC) score to selected input and output hidden states. We evaluate six open-source language models on four QA benchmarks covering faithfulness and factuality. Across the main benchmark, HIDE achieves approximately 29% mean relative AUC improvement over the better of the perplexity and energy baselines in each model/dataset setting, and is competitive with multi-pass methods. Under standard HuggingFace inference, its mean latency reduction relative to five-generation EigenScore is 50.8%. Timing measurements extending to Gemma-2-27B show that overhead varies across models and datasets. A separate comparison on the same 2,000 answers per setting finds that attention mass performs better in seven of eight settings, while HIDE performs better on Gemma-2-9B/RACE. A control based only on the number of selected tokens closely tracks HIDE, limiting the interpretation of its performance as evidence of nonlinear dependence. These results establish the strengths and limitations of HIDE as a detector that uses hidden states without requiring attention weights.*¹

1. Introduction

Our evolving relationship with *truth* in the age of artificial intelligence is increasingly shaped by modern Language Models (LMs) (Brown et al. 2020; Chowdhery et al. 2023; Touvron et al. 2023). Despite their remarkable fluency, these models can generate content that is factually incorrect or unfaithful to the input context while remaining highly confident — a phenomenon commonly referred to as *hallucination* (Ji et al. 2023). Such hallucinated outputs, which often appear coherent and convincing, pose serious challenges to the reliable and trustworthy deployment of LMs, particularly as these systems are integrated into real-world applications at scale (Zhang et al. 2023; Weidinger et al. 2022; Huang et al. 2025; Maynez et al. 2020).

Hallucinations manifest in different forms. Huang et al. (2025) distinguish between *faithfulness hallucinations*, where the output contradicts the provided source or context, and *factuality hallucinations*, where the model generates information inconsistent with real-world facts. Prior work has attributed these failures to factors such as training data artifacts, model architectures, and inference strategies (Lee et al. 2022; Onoe et al. 2022;

* Corresponding author: anwoychatterjee@gmail.com

1 Code and retained experimental records: <https://github.com/C-anwoy/HIDE>.




Figure 1: Our proposed method, HIDE, is a single-pass, training-free approach for hallucination detection. For a given input prompt, HIDE extracts representative tokens from both the input and the LM-generated output, computes a HIDE score using an adapted variant of unbiased HSIC on the corresponding hidden-state embeddings from a specific LM layer, and compares this score against a threshold to identify hallucinated generations.

Zheng, Huang, and Chang 2023). While the term “hallucination” itself is sometimes debated due to its anthropomorphic connotations (Maleki, Padmanabhan, and Dutta 2024), the underlying issue of unreliable generation remains a critical concern, motivating the development of effective hallucination detection methods.

A common approach to hallucination detection estimates model uncertainty using metrics such as perplexity (Ren et al. 2023), entropy (Malinin and Gales 2021), or energy-based scores (Liu et al. 2020), under the assumption that hallucinations correlate with low confidence. However, LMs are often confidently wrong, limiting the effectiveness of uncertainty-based signals (Simhi et al. 2025). Another line of work detects hallucinations via consistency checks, generating multiple outputs per input and comparing them using lexical or semantic similarity measures (Lin, Liu, and Shang 2022; Manakul, Liusie, and Gales 2023). While effective, these approaches incur substantial computational overhead due to repeated auto-regressive decoding. Recent mechanistic studies further show that hallucination-related cues are encoded within LM hidden states (Azaria and Mitchell 2023; Ji et al. 2024; Duan, Yang, and Tam 2024), and some methods train probing classifiers on these representations. In contrast, our goal is to directly leverage internal representations for hallucination detection *without* supervised training and *without* requiring multiple generations.

To this end, we propose **HIDE** (Hallucination DetectIon via Decoupled REpresentations), a training-free method motivated by the hypothesis that weaker dependence between input and output representations can indicate hallucination. As illustrated in Figure 1, HIDE extracts hidden states corresponding to key tokens from both input and output sequences and applies an adapted variant of the Hilbert–Schmidt Independence Criterion (HSIC) to quantify representational coupling. We test whether

lower values of this adapted score help identify incorrect outputs. Unlike multi-pass methods, HIDE requires only a single auto-regressive generation per input, avoiding additional sampled responses; the resulting latency depends on the inference implementation.

We evaluate HIDE on four question answering benchmarks — SQuAD (Rajpurkar et al. 2016) and RACE (Lai et al. 2017) for faithfulness hallucinations, and Natural Questions (Kwiatkowski et al. 2019) and TriviaQA (Joshi et al. 2017) for factuality hallucinations — across six open-source LMs of varying scales and training paradigms, including base and instruction-tuned variants of Llama-3.2-3B, Llama-3-8B (Grattafiori et al. 2024), and Gemma-2-9B (Team et al. 2024). These QA settings allow objective verification of hallucinations against reference answers; we additionally include dialogue experiments to assess generalization beyond QA.

Our main contributions are:

- We introduce a training-free detector computed from selected hidden states during one generation, without attention-weight extraction, and evaluate it on faithfulness and closed-book factuality benchmarks.
- We evaluate six models across four QA datasets, compare with a supervised probe, and examine dialogue, decoding, and other design choices. A direct comparison with attention mass and input-derived norm baselines shows where each detector performs well. A selected-token-count control helps assess what the HIDE score measures.
- We measure a 50.8% mean latency reduction relative to five-generation EigenScore under standard HuggingFace inference. Separate measurements of the overhead added to generation extend to Gemma-2-27B and show how it varies across models and datasets.

The ablations examine layer, kernel, and token-selection choices (§7). We discuss these results alongside the attention and token-count comparisons to distinguish detection performance from evidence about the proposed mechanism. Together, these experiments show when HIDE is useful and where simpler signals perform as well or better.

2. Background and Related Work

Though hallucination is often deemed to be essential for the creativity of LMs (Lee 2023; Chakraborty and Masud 2024), it remains one of the most frequently criticized shortcomings of these models. While some studies argue that hallucinations are an inherent consequence of their stochastic nature, often known as the ‘stochastic parrots’ phenomenon (Bender et al. 2021), many others attribute the issue primarily to artifacts in the data that these models ingest. There is a substantial body of research addressing diverse aspects of hallucination in LMs – we briefly summarize a few of these works related to two major dimensions: sources of hallucination and hallucination detection.

2.1 Sources of Hallucination

Hallucinations in LMs have been attributed to a complex interplay of factors spanning data quality, model architecture, and inference processes. On the data front, pre-training on vast and often imperfect corpora introduces risks of propagating misinformation, factual inaccuracies, and social biases (Lin, Hilton, and Evans 2022; Lee et al. 2022; Hernandez et al. 2022; Paullada et al. 2021; Narayanan Venkit et al. 2023; Ladhak et al. 2023; Kasai et al. 2023). The ever-increasing scale of training data, while enabling greater

generalization, can also amplify errors and result in LMs memorizing, rather than reasoning about information.

At the level of model training and design, certain architectural limitations – such as insufficient bidirectional representation or attention mechanism failures – can introduce systematic vulnerabilities to hallucination (Li et al. 2023; Liu et al. 2023; Wang and Sennrich 2020; Sharma et al. 2024). Moreover, the alignment process that tunes LMs to follow human preferences may occasionally introduce a mismatch between model capabilities and desired behaviors, leading to unexpected model outputs.

Inference strategies, especially those introducing stochasticity or prioritizing diversity (such as top- k or nucleus sampling), have also been linked to hallucination (Fan, Lewis, and Dauphin 2018; Holtzman et al. 2020; Dziri et al. 2021; Chuang et al. 2024b). Additional challenges arise from limitations in context integration and representational bottlenecks in output layers (Liu et al. 2024; Shi et al. 2024; Yang et al. 2018; Chang and McCallum 2022). Together, these findings indicate that hallucinations arise not from a single source but from interactions among data, architecture, and inference.

2.2 Hallucination Detection

Given the broad impact of hallucination, a wide range of detection strategies has been proposed. Many early efforts focused on evaluating model confidence, with the intuition that LMs are less certain when hallucinating. Metrics such as perplexity, entropy, and energy-based scores (Ren et al. 2023; Malinin and Gales 2021; Liu et al. 2020) have been explored for this purpose. Another family of approaches leverages output consistency: by sampling multiple generations for the same prompt, one can measure lexical or semantic agreement across responses (Lin, Hilton, and Evans 2022; Durmus, He, and Diab 2020; Scialom et al. 2021). High variability is often taken as a proxy for hallucination, and techniques such as BERTScore, NLI-based classification, and question-answering-based evaluation have all been applied (Falke et al. 2019; Maynez et al. 2020; Nan et al. 2021; Chen et al. 2024b). However, these multi-pass strategies are computationally demanding and poorly suited to low-latency settings, as they require generating and analyzing numerous outputs for each input.

Other detection frameworks rely on verifying factuality by grounding model outputs in external knowledge sources (Min et al. 2023; Huo, Arabzadeh, and Clarke 2023; Chern et al. 2023; Chen et al. 2024b). While these methods can be effective, they require access to curated external databases and ongoing maintenance, and may not generalize across domains.

More recently, studies have explored the use of LMs themselves as evaluators – either via prompt engineering that elicits self-critique, or by prompting additional LMs to judge the factuality or faithfulness of generated content (Xiong et al. 2024; Kadavath et al. 2022; Manakul, Liusie, and Gales 2023; Chiang and Lee 2023; Luo, Xie, and Ananiadou 2023; Laban et al. 2023; Miao, Teh, and Rainforth 2023). These meta-evaluation techniques provide flexibility in black-box settings, but can inherit the blind spots of the underlying models and are themselves subject to hallucination.

Mechanistic Interpretability and Representation-based Hallucination Detection. A growing body of work in mechanistic interpretability investigates how hallucinations are reflected in the internal representations of language models. Several studies demonstrate that hidden states encode signals predictive of factuality or truthfulness, and leverage supervised probes to extract these signals for hallucination detection or analysis (Azaria and Mitchell 2023; Ji et al. 2024; Duan, Yang, and Tam 2024; Obeso et al. 2025). Recent

work has also explored latent-space diversity, activation geometry, and representation dynamics as indicators of hallucinated generation (Chen et al. 2024a; Geva et al. 2023). Collectively, these studies provide strong evidence that hallucinations are systematically encoded within LM hidden states. While effective, most representation-based detection methods rely on supervised training of task- and model-specific classifiers, or require dataset-dependent calibration. As a result, their performance can degrade under domain or task shifts, and their deployment typically requires access to labeled hallucination data and retraining for new settings. In contrast, HIDE does not train any auxiliary model and instead directly measures statistical dependence between input and output representations at inference time. This design makes HIDE training-free and task-agnostic, enabling deployment without labeled hallucination data. In Section 6, we empirically compare HIDE against a strong supervised probing baseline and show that, while probes perform well on their training domain, HIDE exhibits more robust generalization to unseen datasets without any training.

2.3 Background of Hallucination Detection Measures

Here we briefly discuss the mathematical formulation of a few hallucination detection measures, which are used as baselines in this work. Since our proposed method doesn't involve any training, we compare only against other training-free methods.

Assume that $\mathbf{S}_{in} = \{s_1, s_2, \dots, s_I\}$ is the input prompt, and $\mathbf{S}_{out} = \{t_1, t_2, \dots, t_O\}$ is the corresponding output generated by the LM, parameterized by θ , where $\mathbf{S}_{out}$ is a sequence of O tokens. In case of measures requiring multiple generations for the same input $\mathbf{S}_{in}$, assume $\mathcal{S} = [\mathbf{S}_{out}^1, \mathbf{S}_{out}^2, \dots, \mathbf{S}_{out}^N]$ to be the collection of N generations, where diversity is ensured through stochastic sampling strategies, like top- k , nucleus or temperature sampling. Also, $\mathbb{P}_{LM}(\cdot)$ denotes the probability distribution generated by the LM for any position in the sequence.

Perplexity. Perplexity, as a token-level uncertainty measure, can be utilized for evaluating potential hallucinations (Ren et al. 2023). Perplexity is monotonically related to the mean negative log-likelihood (MNLL) of the output sequence, and hence MNLL can be utilized as a tractable proxy for perplexity as a hallucination detection measure (Chen et al. 2024a). Formally, MNLL can be defined as: $\text{MNLL}(\mathbf{S}_{out}|\mathbf{S}_{in};\theta) = -\frac{1}{O} \sum_{i=1}^O \log \mathbb{P}_{LM}(t_i|t_{<i}, \mathbf{S}_{in})$. A lower value of this measure indicates that the model is more confident about its generation.

Energy Score. The Energy Score (Liu et al. 2020) approaches hallucination detection from a different perspective by leveraging energy-based models. Unlike softmax-based confidence scores, energy scores are aligned with the probability density of the inputs. Let $\mathbb{L}_{LM}(\cdot)$ denote the logit distribution over the set of tokens in the vocabulary $\mathcal{V}$, generated by the LM, such that $\mathbb{P}_{LM}(\cdot) = \text{softmax}(\mathbb{L}_{LM}(\cdot))$. The energy score, for a temperature parameter T , can then be defined as: $\text{Energy}(\mathbf{S}_{in};\theta) = -T \cdot \log \sum_{t \in \mathcal{V}} e^{\mathbb{L}_{LM}(t|\mathbf{S}_{in})/T}$. The energy score is generally used for out-of-distribution detection, where samples with higher energies can be interpreted as datapoints with a lower likelihood of occurrence.

Length-Normalized Entropy. Unlike perplexity which utilizes a single generation for quantifying sequence-level uncertainty, length-normalized entropy (LN-Entropy) uses multiple generations to measure the degree of hallucination (Malinin and Gales 2021). LN-Entropy is defined as: $\text{LNE}(\mathcal{S}|\mathbf{S}_{in};\theta) = -\mathbb{E}_{\mathbf{S}_{out} \in \mathcal{S}} \left[\frac{1}{|\mathbf{S}_{out}|} \sum_{i=1}^{|\mathbf{S}_{out}|} \log \mathbb{P}_{LM}(t_i|t_{<i}, \mathbf{S}_{in}) \right]$. It

Method	Uses	Checks	No. of Generations
	Uncertainty	Consistency	Required
Perplexity	✓	✗	1
Energy	✓	✗	1
Length-Normalized Entropy	✓	✗	N
Lexical Similarity	✗	✓	N
Eigenscore	✗	✓	N
HIDE (Ours)	✗	✗	1

Table 1: Comparison of some popular training-free hallucination detection techniques. Most methods rely on either model uncertainty or consistency across multiple generations (N) – approaches relying on the latter show increased latency and compute cost. HIDE, on the other hand, requires only a single generation and leverages internal representation independence, offering a lightweight yet effective alternative.

is assumed that when an LM generates hallucinated content, it is typically uncertain, resulting in an output distribution with higher entropy (Kadavath et al. 2022).

Lexical Similarity. Lexical Similarity (LS) (Lin, Liu, and Shang 2022) aims to quantify the similarities between the multiple generations from an LM corresponding to the same input prompt, with the assumption that LS will be low in case of hallucinations. It can formally defined as: $LS(\mathcal{S}|\mathbf{S}_{in}; \theta) = \frac{2}{N(N-1)} \sum_{i=1}^N \sum_{j=i+1}^N sim(\mathbf{S}_{out}^i, \mathbf{S}_{out}^j)$, where, $sim(\cdot, \cdot)$ can be any function to quantify similarity, like Rouge-L (Lin 2004).

EigenScore. EigenScore (Chen et al. 2024a) leverages the hidden states of LMs for hallucination detection. It measures semantic diversity in the latent space of LMs’ internal representations by calculating the log-determinant of the covariance matrix of multiple output representations: $ES(\mathcal{S}|\mathbf{S}_{in}; \theta) = \frac{1}{N} \log \det(\Sigma + \alpha \cdot \mathbb{I}_N) = \frac{1}{N} \sum_{i=1}^N \log(\lambda_i)$, where, the covariance matrix of the output representations is given by: $\Sigma = Z^T \cdot C_d \cdot Z$ with $Z = [z_1, z_2, \dots, z_N] \in \mathbb{R}^{d \times N}$ being the embedding matrix of the N generated outputs and $C_d = \mathbb{I}_d - \frac{1}{d} \mathbf{1}_d \mathbf{1}_d^T$ being the centering matrix. The set of eigenvalues of the regularized covariance matrix $\Sigma + \alpha \cdot \mathbb{I}$ is given by $\lambda = \{\lambda_1, \lambda_2, \dots, \lambda_N\}$. It is assumed that when the LM is not hallucinating, the representations of the outputs will be highly correlated, thereby having $\lambda_i = 0$ for most i ’s.

Table 1 presents a summary of the characteristics of the discussed hallucination detection measures, and compares our proposed method with these techniques.

3. Theoretical Foundations

HIDE is motivated by the hypothesis that hallucinated outputs may be less strongly coupled to the input in the model’s hidden representations. We use HSIC to study this possibility. HSIC measures statistical dependence through kernels, which can capture nonlinear relationships that a geometric comparison such as cosine similarity may miss. We distinguish these properties of HSIC from the behavior of the adapted score used in HIDE.

3.1 Measuring Representational Dependence

HSIC measures dependence between two random variables in a reproducing kernel Hilbert space, where a kernel defines similarities between observations. With suitable kernels, population HSIC is zero if and only if the variables are independent. Appendix A.1 gives the formal definition and assumptions. These properties motivate the method; they do not establish that a low HIDE score proves independence or that independence causes hallucination.

3.2 Empirical Estimation

In practice, HSIC is estimated from a finite sample. The unbiased estimator described in Appendix A.2 requires at least four observations. Because selected input or output token sets can be smaller, HIDE adapts that estimator to handle short sequences (Section 4.2). This adaptation is biased at finite sample sizes. Its detection performance and its relation to simpler scores are evaluated separately in Sections 6 and 7.

4. HIDE: Hallucination Detection via Decoupled Representations

Building upon the theoretical foundations of HSIC, as introduced in §3, we propose HIDE, a single-pass approach for detecting hallucinations in LMs by investigating the internal model representations for input and output. HIDE is based on the premise that hallucinations manifest as a quantifiable statistical decoupling between the LM’s internal representations of the input sequence and the generated output. We hypothesize that:

- (i) **Sound Generation:** When an LM generates content that is grounded in the input context or in its parametric knowledge, its internal representation of the output, $H_{out}^{(\ell)}$, remains strongly statistically dependent on the internal representation of the input, $H_{in}^{(\ell)}$, at a given layer ℓ . This strong coupling should result in a substantially positive HSIC value.
- (ii) **Hallucinated Generation:** Conversely, when an LM hallucinates, the generative process decouples the output representation $H_{out}^{(\ell)}$ from the input representation $H_{in}^{(\ell)}$. As the generated content diverges, these representations become increasingly independent. In such cases, their joint distribution approaches the product of their marginals, leading to an HSIC value close to zero, i.e., $\text{HSIC}(H_{in}^{(\ell)}, H_{out}^{(\ell)}) \approx 0$.

These statements motivate our use of HSIC. The implemented method uses the adapted finite-sample score defined below and treats lower values as a signal of possible hallucination. Neither this score nor its correlation with correctness directly tests the independence of input and output representations.

4.1 Problem Formulation

As defined previously in §2.3, $\mathbf{S}_{in} = \{s_1, s_2, \dots, s_I\}$ is an input token sequence of length I , and $\mathbf{S}_{out} = \{t_1, t_2, \dots, t_O\}$ is the corresponding LM-generated output sequence of length O . For a chosen layer ℓ , the matrix of hidden states for the input sequence is $\mathbf{H}_{in}^{(\ell)} \in \mathbb{R}^{I \times d}$, where $\mathbf{h}_{i,in}^{(\ell)} \in \mathbb{R}^d$ is the hidden state for input token s_i , and d is the

hidden state dimensionality. Similarly, for the output sequence, $\mathbf{H}_{out}^{(\ell)} \in \mathbb{R}^{O \times d}$, with rows $\mathbf{h}_{j,out}^{(\ell)} \in \mathbb{R}^d$.

To apply HSIC, we define random variables X and Y based on these representations:

Definition 1

Let $\mathcal{T}_{in}$ be the set of all possible token types that can be selected from an input sequence, and $\mathcal{T}_{out}$ be the set of all possible token types from an output sequence. We define two random variables, X and Y , as :

- $X : \mathcal{T}_{in} \rightarrow \mathbb{R}^d$, where for any selected input token type $\sigma_{in} \in \mathcal{T}_{in}$ (which, in a specific instance, corresponds to a token $s_i \in \mathbf{S}_{in}$), $X(\sigma_{in}) = \mathbf{h}_{i,in}^{(\ell)}$. The codomain of X is $\mathcal{X} = \mathbb{R}^d$.
- $Y : \mathcal{T}_{out} \rightarrow \mathbb{R}^d$, where for any selected output token type $\sigma_{out} \in \mathcal{T}_{out}$ (which, in a specific instance, corresponds to a token $t_j \in \mathbf{S}_{out}$), $Y(\sigma_{out}) = \mathbf{h}_{j,out}^{(\ell)}$. The codomain of Y is $\mathcal{Y} = \mathbb{R}^d$.

For a single input-output pair $(\mathbf{S}_{in}, \mathbf{S}_{out})$, we independently select n_{eff} tokens from $\mathbf{S}_{in}$ and $\mathbf{S}_{out}$, respectively, to obtain samples $\{x_1, \dots, x_{n_{eff}}\}$ for X and $\{y_1, \dots, y_{n_{eff}}\}$ for Y . These form the two sets of observations for calculating empirical HSIC.

4.2 HIDE Score Computation

The HIDE framework for detecting hallucination in a single generation pass is summarized in Algorithm 1. By *single-pass*, we mean that HIDE requires no additional forward passes of the language model beyond the standard auto-regressive generation of a single output sequence. The keyword extraction step operates on the generated text and uses a lightweight embedding-based ranking model; it does not alter the LM's generation process or require additional decoding steps. Consequently, the dominant computational cost of HIDE remains the original generation pass, distinguishing it from multi-pass detection methods that require multiple generations per input.

A key component of the proposed method is the computation of our HSIC-based score. The standard unbiased HSIC estimator, $\widehat{\text{HSIC}}_u$ (c.f. Equation A.3), provides desirable statistical properties but is formally defined for $n \geq 4$ samples. In practical scenarios, particularly when processing short input or output sequences, the effective number of selected tokens, n_{eff} , may fall below this threshold. To ensure the applicability of our framework across all sequence lengths and maintain numerical stability, we employ a pragmatically adapted HSIC-like score, denoted by $\widehat{\text{HSIC}}_{\text{HIDE}}$. This adaptation modifies the denominators in the standard unbiased HSIC estimator (Equation A.3) by replacing terms of the form $(n - c)$ with n , and substituting the overall scaling factor $1/(n(n - 3))$ with $1/n^2$.

Definition 2

For an effective sample size $n \geq 1$, the $\widehat{\text{HSIC}}_{\text{HIDE}}$ score is computed as:

$$\widehat{\text{HSIC}}_{\text{HIDE}}(X, Y) = \frac{1}{n^2} \left[\text{Tr}(\tilde{\mathbf{K}}_X \tilde{\mathbf{K}}_Y) + \frac{(\mathbf{1}^T \tilde{\mathbf{K}}_X \mathbf{1})(\mathbf{1}^T \tilde{\mathbf{K}}_Y \mathbf{1})}{n^2} - \frac{2}{n} \mathbf{1}^T \tilde{\mathbf{K}}_X \tilde{\mathbf{K}}_Y \mathbf{1} \right] \quad (1)$$

where $\tilde{\mathbf{K}}_X, \tilde{\mathbf{K}}_Y$ are the $n \times n$ Gram matrices with zeroed diagonals, and $\mathbf{1}$ is an $n \times 1$ vector of ones.

While this adaptation introduces bias relative to the standard $\widehat{\text{HSIC}}_u$, it yields a score that is well-defined and numerically stable for any $n \geq 1$. This practical modification allows for consistent application of our method across all encountered data instances. We use this adapted score as a detection heuristic and evaluate it in §6. It need not be zero for constant representations at finite n : when both Gram matrices contain only ones before their diagonals are removed, its value is $(n-1)/n^2$. We therefore also compare it with this selected-token-count control in Appendix F.

Algorithm 1 HIDE: Hallucination Detection via Decoupled Representations

Input: Prompt $\mathbf{S}_{in}$, model M , layer ℓ , token budget k , RBF parameter γ , threshold τ
Output: Hallucination or Non-hallucination

- 1: **procedure** DETECTHALLUCINATION($\mathbf{S}_{in}, M, \ell, k, \gamma, \tau$)
- 2: Generate one output $\mathbf{S}_{out}$, caching layer- ℓ hidden states
- 3: Store prompt states in $\mathbf{H}_{in}^{(\ell)}$ and available output states in $\mathbf{H}_{out}^{(\ell)}$
- 4: Exclude the final output token, which has no cached state
- 5: **if** no output state is available **then**
- 6: score $\leftarrow 0$
- 7: **else**
- 8: Rank up to k keywords in the prompt and available output text with KeyBERT
- 9: Map keywords to token positions, retaining the rank order and repeated matches
- 10: If either position list is empty, use its first available tokens, up to k
- 11: $n \leftarrow \min(k, \text{length of input list}, \text{length of output list})$
- 12: Keep the first n positions from each list and collect their states as $\mathbf{X}, \mathbf{Y}$
- 13: Build RBF Gram matrices from $\mathbf{X}$ and $\mathbf{Y}$ using γ
- 14: score $\leftarrow \widehat{\text{HSIC}}_{\text{HIDE}}(\mathbf{X}, \mathbf{Y})$, using Equation 1
- 15: **end if**
- 16: **if** score $< \tau$ **then**
- 17: **return** Hallucination
- 18: **else**
- 19: **return** Non-hallucination
- 20: **end if**
- 21: **end procedure**

Theoretical Properties. Appendix A.3 describes the adapted score for small sample sizes and its limiting behavior under independent and identically distributed paired observations. Selected tokens within one answer need not meet that sampling assumption. The theoretical result therefore does not establish what drives the score’s detection performance at the token budgets used here.

4.3 Methodological Choices

Our design involves two key choices that must be fixed before the HIDE score can be evaluated: (i) how to align the *token sets* drawn from the input-output sequences so

that the *same* number of samples are drawn for both random variables X and Y , and (ii) from which *decoder layer*, $\ell \in \{1, \dots, L\}$ the hidden-state matrices $\mathbf{H}_{\text{in}}^{(\ell)}$ and $\mathbf{H}_{\text{out}}^{(\ell)}$ are taken. We formalise the two choices below and motivate the instantiations that are used throughout the paper.

4.3.1 Strategy for Aligning the Sets of Input and Output Samples. Let $\mathbf{H}_{\text{in}}^{(\ell)} \in \mathbb{R}^{I \times d}$ and $\mathbf{H}_{\text{out}}^{(\ell)} \in \mathbb{R}^{O \times d}$ be the matrices of hidden states for the input and output sequences, respectively, at layer ℓ , where I and O may differ by in magnitude. To guarantee that X and Y are constructed from the *same* effective sample size, n_{eff} , we consider two principled strategies:

- (i) **Keyword-based Subsampling.** We set a token budget $k = 20$ unless stated otherwise. KeyBERT ranks up to k keywords from each sequence using maximal marginal relevance, which balances relevance with diversity. We map each keyword, with and without a leading space, to matching token positions. Matches are kept in keyword order, including repeated positions. If no keyword matches are found, we use the first available token positions, up to k . We then truncate both position lists to the smaller of their lengths and k , giving the effective count n_{eff} . The corresponding hidden states form $\mathbf{X}$ and $\mathbf{Y}$. Thus n_{eff} counts selected entries, which need not be distinct tokens. Output selection uses only tokens with cached hidden states, as specified in Algorithm 1. The effect of the token budget is examined in §7.2.
- (ii) **Rank-truncated SVD Alignment.** Alternatively, we avoid lexical heuristics altogether and enforce equal sample size by projecting *both* matrices $\mathbf{H}_{\text{in}}^{(\ell)}$ and $\mathbf{H}_{\text{out}}^{(\ell)}$ to subspaces with the same rank (i.e., n_{eff}). Concretely, for $\mathbf{H}_{\text{in}}^{(\ell)}$ we compute a thin SVD $\mathbf{H}_{\text{in}}^{(\ell)} = \mathbf{U}_X \Sigma_X \mathbf{V}_X^\top$ and retain the n_{eff} largest singular directions,

$$\mathbf{X}_r = \Sigma_X^{(n_{\text{eff}})} (\mathbf{V}_X^{(n_{\text{eff}})})^\top \in \mathbb{R}^{n_{\text{eff}} \times d}. \quad (2)$$

An analogous projection $\mathbf{Y}_r$ is obtained for $\mathbf{H}_{\text{out}}^{(\ell)}$, yielding aligned sample sets $\mathcal{X} = \{\mathbf{x}_r^{(1)}, \dots, \mathbf{x}_r^{(n_{\text{eff}})}\}$ and $\mathcal{Y} = \{\mathbf{y}_r^{(1)}, \dots, \mathbf{y}_r^{(n_{\text{eff}})}\}$. As both $\mathbf{X}_r$ and $\mathbf{Y}_r$ now contain exactly n_{eff} *orthogonal* rows, the empirical Gram matrices in Equation (1) have identical dimensions, eliminating the need for any further padding or truncation. We discuss the performances of HIDE with both these alignment strategies in §7.6. Unless stated otherwise, we adopt the keyword-based subsampling strategy for our experiments.

4.3.2 Layer Selection. Each decoder layer captures different levels of abstraction; therefore, the choice of ℓ can, in principle, influence the measured dependence. Based on prior work suggesting that representations corresponding to the middle layers are often richer (Skean et al. 2025) and provide stronger sentence-level semantic information suitable for detecting inconsistencies (Azaria and Mitchell 2023), we fix $\ell = \ell_{\text{mid}}$ for all our experiments to avoid an additional hyperparameter sweep. However, our framework is flexible to any layer choice and we observe the performance of HIDE to be almost agnostic to the choice of ℓ (§7.3).

4.4 Structural Analysis of Decoupling

We compare HIDE with two quantities that describe how generated tokens use the prompt: attention mass and an input-derived norm. These comparisons test whether

the scores vary together. They do not establish a causal explanation of hallucination. In closed-book QA, an answer can attend to the question and still be factually wrong.

At the selected layer ℓ , let $\alpha_{j,i}^{(h)}$ denote the causal self-attention weight from an observed generated-token position j to prompt token i in head h . Let J be the number of generated-token positions with available cached states. We use

$$\bar{\alpha}_{j,i} = H^{-1} \sum_{h=1}^H \alpha_{j,i}^{(h)}, \quad (3)$$

$$\Omega = J^{-1} \sum_{j=1}^J \sum_{i=1}^I \bar{\alpha}_{j,i}, \quad (4)$$

$$\tilde{\Delta}_{\text{in}} = J^{-1} \sum_{j=1}^J \left\| \sum_{i=1}^I \bar{\alpha}_{j,i} h_{i,\text{in}}^{(\ell-1)} \right\|_2. \quad (5)$$

Here, H is the number of attention heads, Ω is the mean attention mass assigned to prompt tokens, and $\tilde{\Delta}_{\text{in}}$ is the mean norm of attention-weighted input hidden states. The norm omits the learned value and output projections, so it is a proxy for the input-derived update, not the actual update to the residual stream. We use attention weights from the layer used by HIDE and hidden states entering that layer. Both baselines average over all available generated-token positions and require no keyword extraction. The final generated token has no cached hidden state because it is not passed through the model again. If no output state is available, all three scores are set to zero. For detection, we use higher raw scores to predict correctness, without choosing the score direction from test results.

Figure 2 uses the same 2,000 examples per model and dataset as the direct baseline comparison in Section 5. For Llama-3-8B, $\text{PCC}(\text{HIDE}, \Omega)$ is 0.560 on SQuAD and 0.648 on NQ; $\text{PCC}(\text{HIDE}, \tilde{\Delta}_{\text{in}})$ is 0.003 and 0.367, respectively. For Gemma-2-9B, the corresponding attention correlations are 0.565 and 0.430, and the norm correlations are positive: 0.126 on SQuAD and 0.215 on NQ. Thus the relationship between HIDE and the norm varies with model and dataset. These results do not support a general inverse relationship between update magnitude and the precision of information used in an answer.

Question-conditioned but wrong answers. For NQ, star markers identify question-conditioned but wrong (QCBW) answers: incorrect answers in the top quarter of question-answer embedding similarity among incorrect answers, including ties at the cutoff. This subset helps examine answers that remain close to the question in topic but are incorrect. It is an observational comparison, not a causal intervention. All examples remain in the figure and in the reported correlations. Section 6.2 tests the three scores directly as detectors.

5. Experimental Setup

Evaluation Datasets. Following the setup of Hong et al. (2024), we evaluate HIDE on four benchmark QA-datasets, as discussed below, to assess its effectiveness in detecting both faithfulness and factuality hallucinations:




Figure 2: HIDE, prompt attention mass Ω , and input-derived norm proxy $\tilde{\Delta}_{in}$ for Llama-3-8B and Gemma-2-9B. Each model/dataset uses 2,000 seed-42 examples and greedy generation stopped after the first answer line. Columns show SQuAD and NQ; rows compare HIDE with attention mass and the norm, respectively. Answers are correct if answer–reference sentence similarity exceeds 0.9. Teal squares denote incorrect answers, red circles correct answers, and gold stars the NQ QCBW subset: incorrect answers in the top quarter of question–answer similarity, including cutoff ties. PCC uses all examples and is computed before display rounding.

- (i) **RACE (ReAding Comprehension dataset from Examinations)** (Lai et al. 2017) is a multiple-choice reading comprehension dataset from English exams for Chinese students, featuring 27,933 passages and 97,867 questions. We use the test split of the ‘high’ subset containing 3,498 questions. RACE is challenging due to its focus on inference and reasoning. RACE is used to evaluate *faithfulness hallucinations*, as answers must be supported by explicit information in the provided passage, making inference and grounding critical.
- (ii) **SQuAD (Stanford Question Answering Dataset)** (Rajpurkar et al. 2016) is a widely used benchmark for extractive question answering, containing crowd-sourced questions on Wikipedia articles. SQuAD-v1.1 contains only questions for which answers are present within the provided context, while SQuAD-v2.0 extends it to include both answerable and unanswerable questions. Following Chen et al. (2024a), we use only answerable questions from the development split of v2.0, retaining 5,928 context-question-answer pairs. SQuAD also targets *faithfulness hallucinations*, as the correct answer must be a text span from the passage, requiring the model to remain faithful to the context.
- (iii) **NQ (Natural Questions)** (Kwiatkowski et al. 2019) is a large-scale question answering dataset consisting of real, anonymized queries issued to the Google search engine. Following Lin, Trivedi, and Sun (2024), we use the validation set consisting of 3610 pairs. NQ is notable for its use of naturally occurring user queries, making it a realistic and challenging benchmark for end-to-end QA. NQ is primarily used to assess *factuality hallucinations*, as it requires models to provide correct answers to open-domain questions based on external world knowledge, testing factual correctness.

- (iv) **TriviaQA** (Joshi et al. 2017) is a large-scale reading comprehension dataset with over 95,000 crowd-authored question-answer pairs and associated evidence documents. The questions are naturally complex, often requiring cross-sentence reasoning. We use the `rc.nocontext` validation subset, which contains 9,960 deduplicated pairs. TriviaQA also focuses on *factuality hallucinations*, as we use the subset which has no input context, hence the evaluation relies on accurate, fact-based answers.

Dialogue Evaluation. To assess whether the representational decoupling hypothesis extends beyond QA, we additionally evaluate HIDE on the dialogue split of the HaluEval benchmark. In this setting, hallucinations typically manifest as responses that contradict established dialogue context or introduce unsupported factual claims. Unlike QA, dialogue responses are open-ended, but correctness can still be evaluated based on consistency with the preceding conversational turns. This experiment allows us to test whether HIDE generalizes to conversational generation without relying on multiple generations or external verification.

Models and Generation Configurations. To ensure that our findings are broadly applicable and robust across diverse model scales and variants, we consider six open-source LMs, which include both *base* and *instruction-tuned* variants of Llama-3.2-3B, Llama-3-8B, and Gemma-2-9B. All model checkpoints are sourced directly from the HuggingFace model hub without any additional fine-tuning. For all experiments, we generate outputs in a *zero-shot setting*, using greedy decoding for techniques requiring a single output, and set the values of temperature, top- p and top- k parameters to 0.5, 0.99 and 10, respectively for stochastic sampling in case of methods requiring multiple outputs for the same input prompt. **The main benchmark experiments use an NVIDIA A100-SXM4-80GB GPU. The direct attention/norm comparison and the paired timing measurements use A100 80GB PCIe GPUs, as specified below and in Section 6.4.**

Evaluation Metrics. To comprehensively assess the effectiveness of HIDE and baseline methods in hallucination detection, we employ a variety of evaluation metrics commonly used in the literature.

- **Area Under the ROC Curve (AUC-ROC):** This threshold-independent metric quantifies how well a given method can distinguish between hallucinated and non-hallucinated outputs. Higher AUC-ROC values indicate stronger binary discrimination.
- **Pearson Correlation Coefficient (PCC):** **PCC measures the linear correlation between a detector’s scores and the reference correctness measure. For sentence similarity and ROUGE-L, this measure is continuous; for exact match, it is binary.**

We report AUC-ROC and PCC values for both *sentence embedding similarity* and *ROUGE-L overlap* as the reference ground-truth measure, denoted in the results as AUC_s , PCC_s (using sentence similarity) and AUC_r , PCC_r (using ROUGE-L). The sentence similarity is calculated as the cosine similarity between the sentence embeddings of the generated output and the reference answer, using the *nli-roberta-large* model (Reimers and Gurevych 2019) – an output is considered correct if its similarity exceeds a fixed threshold set to 0.9. When using the ROUGE-L measure (Lin 2004), outputs with an overlap score of more than 0.5 with the reference are considered to be correct. We also report the AUC-ROC and PCC values for exact match with the ground-truth answer as the correctness measure.

Direct comparison with attention and norm baselines.. We compare HIDE, attention mass, and the input-derived norm on the first 2,000 examples in each dataset’s seed-42 shuffled order, using Llama-3-8B and Gemma-2-9B: 16,000 answers in total. The sample size was chosen to fit the available compute time after earlier runs had been inspected. Examples were not selected by detector score or correctness. Panel A of Table 4 reproduces the full-benchmark HIDE values from Appendix B; Panel B reports the three detectors on the same 2,000 answers per setting.

The benchmark stopping conditions can allow generation to continue into another question, especially on NQ and TriviaQA. For this comparison, we stop at the first line break after the first nonempty answer line, while retaining the end-of-sequence and stop-token conditions and the 256-token cap. Correctness is evaluated on that answer line. The prompts, model checkpoints, layers, greedy decoding, 20-token budget, default RBF kernel, and score formula are unchanged. HIDE uses the available cached output states, excluding the final token, which has not been passed through the model. NQ and TriviaQA are closed-book, open-domain tasks: no retrieved evidence is supplied.

All three detectors score the same generations against the same first reference answer. AUC uses binary correctness labels; PCC_s and PCC_r use continuous sentence similarity and ROUGE-L, respectively. Higher raw scores predict correctness. We do not choose score direction, kernel parameters, or examples to improve test results. We retain answers that reach the token cap, empty answers, and cases assigned zero scores because output states are unavailable. For confidence intervals, we draw 2,000 bootstrap samples of examples with replacement (seed 42), using the same sampled examples for each detector. Each interval applies to one comparison and is not adjusted for multiple comparisons. Per-example records are available with the code. Because the stopping conditions and sample sizes differ, these results cannot be compared directly with the full-benchmark values.

6. Results and Discussions

We first report the full-benchmark comparison with uncertainty-based and multi-pass methods. We then compare HIDE, attention mass, and the input-derived norm on the same answers under the first-answer-line stopping conditions defined in Section 5. The two experiments are reported separately because their stopping conditions and sample sizes differ.

Table 2 summarizes AUC_s , AUC_r , PCC_s , and PCC_r across four datasets and six models for HIDE and five baseline methods. Table B.1 in Appendix B gives the per-dataset results. Table C.1 in Appendix C reports the results using exact match to determine correctness. We also compare with a supervised probe and examine dialogue and computational efficiency.

Comparison with Supervised Probing Methods. We additionally compare HIDE against a supervised probing baseline following the state-of-the-art methodology of Obeso et al. (2025). Specifically, we train a logistic regression classifier on token-level hidden states extracted from TriviaQA, using the same layer selection strategy as the original work, i.e., extracted from the middle layer. During inference, sequence-level hallucination scores are obtained by aggregating token-level predictions via maximum probability, ensuring sensitivity to localized hallucinated spans. Table 3 compares HIDE against a supervised probing baseline across four datasets and four large-scale LMs. As expected, probing performs strongly on its training distribution. For example, on TriviaQA, prob-

Method	AUC _s	AUC _r	PCC _s	PCC _r
Llama-3.2-3B				
Perplexity	51.84	52.13	2.19	10.76
Energy	41.53	41.05	-19.18	-12.51
LN-Entropy	62.86	62.42	15.84	19.52
Lexical Similarity	68.61	70.65	31.04	36.45
Eigenscore	66.98	66.65	25.39	25.85
HIDE	69.37	69.13	35.86	28.67
Llama-3.2-3B-Instruct				
Perplexity	56.75	57.70	10.87	18.77
Energy	39.99	40.48	-19.11	-10.48
LN-Entropy	65.29	65.70	21.77	24.19
Lexical Similarity	71.38	72.75	38.38	40.69
Eigenscore	72.68	72.57	38.44	37.19
HIDE	71.26	71.20	38.68	33.47
Llama-3-8B				
Perplexity	45.56	45.68	-3.92	2.33
Energy	36.63	36.36	-21.55	-18.13
LN-Entropy	60.99	60.87	15.58	18.31
Lexical Similarity	68.12	69.50	31.45	35.90
Eigenscore	66.54	66.81	26.84	28.79
HIDE	72.84	72.60	39.32	35.60
Llama-3-8B-Instruct				
Perplexity	56.71	57.47	6.38	15.13
Energy	42.25	42.53	-13.99	-6.88
LN-Entropy	66.69	66.60	24.15	27.60
Lexical Similarity	70.48	71.89	36.83	39.45
Eigenscore	74.12	73.99	41.57	39.83
HIDE	73.58	73.43	41.86	37.84
Gemma-2-9B				
Perplexity	45.07	47.68	5.97	11.43
Energy	57.49	56.00	2.19	-1.60
LN-Entropy	52.25	55.81	12.70	20.58
Lexical Similarity	63.97	66.21	26.68	33.13
Eigenscore	70.67	71.37	24.37	29.18
HIDE	73.29	72.91	35.70	31.05
Gemma-2-9B-Instruct				
Perplexity	44.25	47.29	25.49	10.43
Energy	76.73	75.20	8.61	23.92
LN-Entropy	53.57	56.06	21.65	18.39
Lexical Similarity	62.43	65.27	31.27	29.45
Eigenscore	67.26	69.08	23.24	28.73
HIDE	73.24	73.49	38.87	35.97

Table 2: Full-benchmark performance under the original generation/evaluation protocol, averaged across the four QA datasets. The new corrected paired comparison is reported separately in Table 4. We report average AUC-ROC and PCC values using sentence similarity (s) and ROUGE-L (r) as correctness measures. Multi-pass methods are highlighted in gray. The best overall result is shown in **bold**, and the best among single-pass methods is underlined. Refer to Table B.1 of Appendix B for detailed breakdown.

Method	RACE				SQuAD				NQ				TriviaQA			
	AUC _s	AUC _r	PCC _s	PCC _r	AUC _s	AUC _r	PCC _s	PCC _r	AUC _s	AUC _r	PCC _s	PCC _r	AUC _s	AUC _r	PCC _s	PCC _r
Llama-3-8B																
Probing	67.25	67.54	26.54	34.60	68.79	69.83	37.06	37.68	75.27	76.07	41.86	32.14	85.18	87.68	77.42	70.81
HIDE	69.81	66.95	16.76	22.87	76.87	77.45	49.09	50.43	79.02	79.17	49.61	33.91	65.65	66.82	41.82	35.19
Llama-3-8B-Instruct																
Probing	70.83	68.16	33.15	37.61	62.88	62.52	27.15	26.34	76.70	77.35	47.20	35.74	84.88	86.95	75.67	66.77
HIDE	73.20	72.73	22.12	30.11	80.15	79.14	52.27	55.98	79.29	78.91	56.33	37.98	61.70	62.94	36.73	27.27
Gemma-2-9B																
Probing	57.92	58.25	15.71	13.47	69.25	70.36	36.51	38.15	85.97	84.32	24.44	16.55	86.70	87.28	47.67	40.62
HIDE	58.06	57.99	20.60	16.50	75.17	76.66	47.63	48.43	79.08	75.96	24.08	15.41	80.84	81.02	50.51	43.88
Gemma-2-9B-Instruct																
Probing	66.66	66.64	20.03	22.70	55.35	56.89	15.29	14.99	84.75	83.83	11.09	22.16	77.08	78.42	-2.06	19.82
HIDE	64.18	64.03	11.92	16.26	72.69	73.32	39.73	39.70	85.19	85.89	51.20	41.99	70.90	70.72	52.64	45.90

Table 3: Comparison between supervised probing and HIDE across four QA datasets. Probing models are trained on TriviaQA and evaluated without retraining on other datasets. Bold indicates the best-performing method per metric.

ing achieves AUC_s scores of 85.18 and 84.88 on Llama-3-8B and Llama-3-8B-Instruct respectively, with correspondingly high PCC values, reflecting the benefit of supervised alignment.

However, this advantage does not consistently transfer across datasets. On SQuAD, HIDE outperforms probing across all four models, with gains exceeding 8 AUC points on Llama-3-8B (76.87 vs. 68.79) and 17.27 points on Llama-3-8B-Instruct (80.15 vs. 62.88). The NQ comparison is mixed: for example, Gemma-2-9B has lower AUC_s with HIDE than with the probe (79.08 versus 85.97), whereas Gemma-2-9B-Instruct has a small HIDE advantage (85.19 versus 84.75).

HIDE requires no labeled examples or detector training. It performs better than the probe on SQuAD, while the NQ comparison is mixed. These results show that it can be useful when labeled hallucination data are unavailable; they do not establish that it is uniformly better under changes in domain.

6.1 Faithfulness Hallucination Detection

On SQuAD and RACE, answers must be supported by the supplied context. HIDE has higher AUC_s than both uncertainty baselines, Perplexity and Energy, in 10 of the 12 model/dataset settings (Table B.1). Its mean relative improvement over the better of these two baselines is 47.4% on SQuAD and 20.1% on RACE.

HIDE’s advantage persists when compared to the state-of-the-art multi-pass methods. On SQuAD, for example, HIDE achieves an AUC_s score of 80.15% with Llama-3-8B-Instruct, which is an absolute improvement of 8 percentage points over Eigenscore, the state-of-the-art multi-pass baseline. The improvements are also reflected in PCC values, with HIDE outperforming single-pass baselines and matching the best multi-pass alternatives in most settings. This trend is also mirrored on the RACE dataset, where HIDE provides consistently robust performance across both base and instruction-tuned variants of all models. Overall, for faithfulness hallucination detection, HIDE achieves an average relative gain of 33.7% and 4.2% in AUC_s over the single-pass and multi-pass baselines, respectively. These benchmark relative gains average the per-model/dataset improvement over the best evaluated uncertainty-based or multi-pass baseline in that cell; the corrected attention comparison below is a separate experiment.

6.2 Factuality Hallucination Detection

We further evaluate HIDE on benchmarks designed to test factuality hallucinations, such as Natural Questions (NQ) and TriviaQA. Here, the outputs may be fluent yet factually unsupported. As shown in Tables B.1 and C.1, even in the case of factuality hallucinations, HIDE has higher AUC_s than both uncertainty baselines in 9 out of 12 settings, with an average relative gain of 23.3% in AUC_s scores across models for NQ and TriviaQA. For instance, on NQ, HIDE achieves AUC_s score of 79.29% with Llama-3-8B-Instruct, compared to 57.47% for the best uncertainty-based single-pass competitor. The PCC comparison varies across models and datasets.

Compared to multi-pass methods, HIDE is generally competitive and, in some settings, outperforms both Eigenscore and Lexical Similarity. For example, on NQ, HIDE has higher AUC_s than the five unsupervised baselines in Appendix B for all six model variants; PCC comparisons are more mixed. However, in certain configurations, such as TriviaQA with Llama-3-8B, the advantage of HIDE is less pronounced, and Eigenscore achieves marginally higher AUC-ROC and PCC values. A closer examination of these cases suggests that datasets with extremely short or repetitive answers can reduce the discriminative power of representation-based dependence measures. We discuss such limitations in detail in §8.




Figure 3: Benchmark latency comparison under standard HuggingFace inference on an A100-SXM4-80GB. Each bar is the equal-weight mean of the per-example latencies for RACE, SQuAD, NQ, and TriviaQA for the indicated model. HIDE uses one generation; multi-pass baselines use five outputs. The HIDE total includes generation, keyword extraction, and scoring. The mean reduction relative to EigenScore is 50.8%, calculated by averaging the six model-specific percentage reductions after averaging datasets within each model. Paired generation-overhead measurements on A100 PCIe are reported in Table 6; different hardware, cohorts, and measurement protocols are not pooled.

Matched comparison with attention and norm baselines. Table 4 reports the direct comparison. Panel A reproduces the full-benchmark HIDE values from Appendix B for reference; Panel B evaluates all three detectors on the same answers under Section 5. Only comparisons within Panel B use matched examples and stopping conditions. Attention has higher AUC_s in seven of eight settings. HIDE performs better on Gemma-2-9B/RACE (57.93 versus 51.23), while attention performs better in all four closed-book factuality settings: Llama/NQ (84.24 versus 78.47), Llama/TriviaQA (73.68 versus 57.01), Gemma/NQ (67.97 versus 62.74), and Gemma/TriviaQA (66.94 versus 49.83). The norm also performs better than HIDE on both TriviaQA settings. Attention to the question is therefore a useful detection signal even without a supplied evidence passage. Appendix F reports ROUGE-L results and confidence intervals for the AUC_s differences.

We also compare HIDE with a score based only on the actual number n of selected tokens: $c(n) = (n - 1)/n^2$ for $n \geq 1$, and zero when no output states are available. This is the value of the adapted score for all-ones Gram matrices before their diagonals are removed. The control closely tracks HIDE (PCC at least 0.9987), with an absolute AUC_s difference below 0.43 percentage points in every setting. Thus the number of selected tokens largely accounts for the score in these runs. This limits attributing its predictive performance specifically to nonlinear dependence between representations.

Interpretation of the baseline comparison. The comparison shows that the choice of detector depends on the setting. On Gemma-2-9B/RACE, HIDE’s AUC_s exceeds attention by 6.69 percentage points (95% confidence interval [3.70, 9.69]); it also performs better using ROUGE-L. Attention performs better in the other seven settings. The full benchmark and the probe, dialogue, and ablation experiments show HIDE’s performance against their respective baselines. Together, these results support HIDE as one detection option, while showing that HSIC is not necessary to obtain a useful signal in every setting.

Panel A: HIDE full-benchmark reference (Appendix B)

Model	Dataset	N	AUC_s	PCC_s
Llama-3-8B	SQuAD	5928	76.87	49.09
Llama-3-8B	RACE	3498	69.81	16.76
Llama-3-8B	NQ	3610	79.02	49.61
Llama-3-8B	TriviaQA	9960	65.65	41.82
Gemma-2-9B	SQuAD	5928	75.17	47.63
Gemma-2-9B	RACE	3498	58.06	20.60
Gemma-2-9B	NQ	3610	79.08	24.08
Gemma-2-9B	TriviaQA	9960	80.84	50.51

Panel B: Matched first-answer-line comparison

Model	Dataset	N_+	HIDE		Attention Ω		Norm $\tilde{\Delta}_{in}$	
			AUC_s	PCC_s	AUC_s	PCC_s	AUC_s	PCC_s
Llama-3-8B	SQuAD	741	76.33	47.83	85.34	58.63	53.79	2.15
Llama-3-8B	RACE	778	61.66	5.35	64.67	9.28	49.97	-15.36
Llama-3-8B	NQ	245	78.47	47.85	84.24	51.08	72.05	37.03
Llama-3-8B	TriviaQA	1153	57.01	13.01	73.68	54.76	64.68	29.55
Gemma-2-9B	SQuAD	679	75.90	48.53	86.72	63.97	54.32	14.68
Gemma-2-9B	RACE	1150	57.93	20.78	51.23	-0.84	50.56	0.18
Gemma-2-9B	NQ	448	62.74	30.06	67.97	46.71	58.63	28.29
Gemma-2-9B	TriviaQA	1257	49.83	-1.76	66.94	39.88	61.53	20.04

Table 4: Panel A reproduces HIDE values from Table B.1 of Appendix B without alteration (N : full-benchmark sample size). Panel B evaluates all three scores on the same 2,000 fixed seed-42 examples per model/dataset under first-answer-line stopping; N_+ refers only to Panel B. Because cohorts and stopping differ, Panel A is a reference, not a matched comparator for Panel B. N_+ counts answers with sentence similarity greater than 0.9. AUC and PCC are multiplied by 100; PCC uses the continuous reference metric. Within Panel B, all scores use the same examples and labels; higher raw scores predict correctness. No direction is selected using test performance.

Attention-based detection is also supported by prior work. Lookback Lens trains a classifier on attention ratios and reports strong detection and transfer results (Chuang et al. 2024a). Both HIDE and the raw attention baseline used here require no detector training. They measure different quantities: attention weights describe how tokens are combined, while the resulting representations also depend on transformed value vectors (Kobayashi et al. 2020). This distinction motivates studying hidden states, but the selected-token-count result prevents us from concluding that HIDE’s advantage comes from additional semantic information.

HIDE can be computed when hidden states are available without extracting attention weights. This distinction can matter in implementations that do not expose those weights. For example, FlashAttention avoids storing the full attention matrix in GPU memory (Dao et al. 2022). Extracting attention mass can require changes to the implementation, although it could be accumulated without storing the full matrix. Hidden-state extraction also has a cost. We therefore identify a difference in required model access, without claiming a measured speed advantage over attention or an advantage in optimized serving.

Method	HaluEval (Dialogue)	
	AUC _s	AUC _r
Llama-3.2-3B		
Perplexity	68.00	<u>72.00</u>
Energy	62.00	66.00
LN-Entropy	68.00	68.00
Lexical Similarity	75.00	72.00
Eigenscore	64.00	61.00
HIDE	<u>70.00</u>	57.00
Llama-3.2-3B-Instruct		
Perplexity	68.00	<u>72.00</u>
Energy	58.00	57.00
LN-Entropy	67.00	69.00
Lexical Similarity	69.00	72.00
Eigenscore	65.00	64.00
HIDE	63.00	59.00
Llama-3-8B		
Perplexity	62.00	68.00
Energy	57.00	63.00
LN-Entropy	61.00	62.00
Lexical Similarity	66.00	68.00
Eigenscore	61.00	58.00
HIDE	66.00	55.00
Llama-3-8B-Instruct		
Perplexity	64.00	70.00
Energy	56.00	62.00
LN-Entropy	65.00	67.00
Lexical Similarity	66.00	70.00
Eigenscore	60.00	59.00
HIDE	61.00	55.00
Gemma-2-9B		
Perplexity	73.00	73.00
Energy	43.00	39.00
LN-Entropy	72.00	71.00
Lexical Similarity	75.00	73.00
Eigenscore	78.00	72.00
HIDE	64.00	53.00
Gemma-2-9B-Instruct		
Perplexity	71.00	73.00
Energy	42.00	43.00
LN-Entropy	68.00	71.00
Lexical Similarity	67.00	72.00
Eigenscore	69.00	71.00
HIDE	56.00	54.00

Table 5: Dialogue hallucination detection results on the HaluEval benchmark. We report AUC-ROC using sentence similarity (AUC_s) and ROUGE-L (AUC_r) as correctness measures. Rows corresponding to multi-pass methods are highlighted in grey. The best result for each metric is shown in **bold**, while the best single-pass result is underlined.

6.3 Generalization to Conversational Hallucinations

To evaluate whether the representational decoupling hypothesis underlying HIDE extends beyond question answering, we examine its performance on the dialogue split of the HaluEval benchmark. In this setting, hallucinations typically manifest as inconsistencies with prior conversational context or unsupported factual claims, rather than incorrect answers to well-defined questions. Dialogue therefore provides a challenging testbed for assessing whether input-output representational dependence remains informative in more open-ended generation scenarios.

Table 5 reports dialogue hallucination detection performance across six language models, including both base and instruction-tuned variants of Llama and Gemma. Overall, HIDE demonstrates competitive and often strong performance across models, particularly on larger architectures. For instance, on Llama-3-8B HIDE achieves the best AUC_s score, substantially outperforming the uncertainty-based single-pass baselines and matching or exceeding multi-pass methods.

At the same time, performance varies across models and metrics. On some smaller or instruction-tuned models, simpler signals such as lexical similarity or energy-based scores perform competitively, and in a few cases outperform HIDE on individual metrics. This variability highlights the increased difficulty of dialogue hallucination

detection, where surface-level similarity or confidence cues may occasionally align with correctness.

These results extend the evaluation beyond QA: HIDE is competitive on some dialogue settings without additional supervision or sampled generations, but simpler baselines perform better on others. They show useful detection performance in some conversational cases, without establishing a general advantage.

6.4 Computational Efficiency

Comparison with multi-generation baselines. Figure 3 reports the six-model, four-dataset comparison under standard HuggingFace inference on an A100-SXM4-80GB. For model m and method b , we first average the four dataset-specific mean latencies equally, obtaining $\bar{T}_{m,b}$. We then report the mean relative reduction

$$R_b = \frac{100}{6} \sum_{m=1}^6 \left(1 - \frac{\bar{T}_{m,\text{HIDE}}}{\bar{T}_{m,b}} \right).$$

The mean latency reductions are 50.8% relative to EigenScore, 50.8% relative to lexical similarity, and 53.2% relative to LN-Entropy. The average across these three comparators is 51.6%. We use the EigenScore comparison when reporting the 50.8% reduction. These measurements compare one-generation and five-generation methods under standard HuggingFace inference.

For Llama-3-8B on SQuAD, base generation takes 1.6169 s and the complete HIDE pipeline takes 1.8445 s, adding 0.2276 s (14.08%). On RACE, base generation takes 1.5021 s and HIDE takes 1.7107 s. All differences and percentages are computed before rounding.

Measuring the overhead added by HIDE. We evaluate Llama-3.2-3B, Llama-3-8B, Gemma-2-9B, and Gemma-2-27B on 200 fixed seed-42 queries from each of SQuAD and NQ. Measurements use standard HuggingFace Transformers/PyTorch, BF16, eager attention, and one A100 80GB PCIe GPU. After ten warm-up base/HIDE pairs, we collect three measured pairs per query, alternate base/HIDE order, synchronize CUDA around measurements, and verify identical generated token IDs within each pair. Model loading and correctness evaluation are outside the measured intervals. Base latency measures generation without hidden-state capture; HIDE total includes generation with hidden-state capture, keyword extraction, score computation, and associated bookkeeping.

These measurements use the benchmark stopping conditions, rather than the first-answer-line conditions of the attention/norm comparison. All eight timing jobs use the same GPU. Other workloads shared the host, and we did not record clock rates or power continuously. The confidence intervals describe variation within the runs; they do not account for all possible effects of host sharing or changes in GPU operating conditions.

Measured overhead. Table 6 and Figure 4 report the paired timing measurements. We average the three repeats within each query, then average over queries. Incremental overhead is $T_{\text{HIDE}} - T_{\text{base}}$; relative overhead is $100(\bar{T}_{\text{HIDE}} - \bar{T}_{\text{base}})/\bar{T}_{\text{base}}$, not the mean of per-query percentages. We bootstrap queries 2,000 times and retain all measurements, including negative paired differences. For Llama-3-8B on SQuAD, base generation is 0.9277 s and HIDE total is 1.1587 s, giving 0.2310 s overhead (24.90%). On NQ, the corresponding values are 0.7906 s, 0.8832 s, and 0.0925 s (11.70%). Differences are com-

Model	Dataset	Base (s)	HIDE total (s)	Overhead (s)	Overhead (%)	95% CI (s)
Llama-3.2-3B	SQuAD	1.2642	1.4642	0.2000	15.82	[0.1741, 0.2279]
Llama-3.2-3B	NQ	0.6215	0.7090	0.0874	14.07	[0.0776, 0.1002]
Llama-3-8B	SQuAD	0.9277	1.1587	0.2310	24.90	[0.1943, 0.2752]
Llama-3-8B	NQ	0.7906	0.8832	0.0925	11.70	[0.0764, 0.1106]
Gemma-2-9B	SQuAD	3.3214	3.5927	0.2713	8.17	[0.2197, 0.3273]
Gemma-2-9B	NQ	1.9465	2.3275	0.3809	19.57	[0.3222, 0.4494]
Gemma-2-27B	SQuAD	4.1989	4.3958	0.1969	4.69	[0.1644, 0.2365]
Gemma-2-27B	NQ	1.1389	1.2195	0.0806	7.08	[0.0751, 0.0864]

Table 6: Paired timing under the benchmark stopping protocol on an A100 80GB PCIe: 200 queries per model/dataset, 10 warm-up pairs, and three measured pairs per query. Total includes generation, hidden-state capture, keyword extraction, and HIDE scoring; model loading and correctness evaluation are excluded. Overhead is total minus base; percentages are ratios of means. All differences are computed before rounding. Intervals use 2,000 query bootstrap draws after averaging repeats; all observations are retained.

puted from unrounded measurements. These paired measurements quantify generation overhead on A100 PCIe; they are not substituted into the SXM4 baseline comparison or used to calculate its 50.8% reduction.

Model scale and complexity. Gemma-2-27B extends the parameter scale substantially: its SQuAD overhead is 0.1969 s (95% CI [0.1644, 0.2365]), and its NQ overhead is 0.0806 s ([0.0751, 0.0864]). These observations support feasibility at the tested scales, not constant absolute overhead across models or datasets. The tested hidden widths range from 3,072 to 4,608; increasing parameter count does not imply proportionately increasing hidden width. Generated lengths also differ across models (Appendix F), preventing a controlled parameter-only scaling inference.

For fixed effective token count n , RBF Gram construction costs $\mathcal{O}(n^2d)$; the implementation’s matrix products add $\mathcal{O}(n^3)$. Thus the score arithmetic is $\mathcal{O}(d)$ when n is fixed. This statement does not cover autoregressive generation, sequence-dependent keyword extraction, hidden-state collection, or total model parameter count. Fixed- n arithmetic complexity does not imply constant end-to-end overhead.

Implications for optimized serving. A single-generation detector avoids requesting additional sampled responses. Its latency in an optimized engine nevertheless depends on batching, access to hidden states or attention weights, and data movement. These factors affect HIDE, the attention/norm baselines, and multi-pass methods that use internal representations. We did not measure vLLM, PagedAttention, or another optimized serving engine. The standard HuggingFace measurements therefore do not establish a speedup, or a lower bound on speedup, in these systems.

6.5 Interpreting the HIDE Score for Hallucination Detection

Under the original decision convention, lower HIDE scores are treated as evidence of a potentially incorrect answer. The threshold analysis below is reported for the benchmark protocol; its operating points have not been validated under the corrected stopping protocol. Given the count-control results in Section 6, a score threshold should not be interpreted as a calibrated measure of semantic dependence. We describe how




Figure 4: Measured incremental HIDE latency versus model parameter count on (a) SQuAD and (b) NQ, using the same paired measurements as Table 6. Each point averages 200 queries with three base/HIDE pairs per query; error bars are pointwise 95% query-bootstrap intervals after averaging repeats. Labels include hidden width d . Measurements use an A100 80GB PCIe, standard HuggingFace inference, and the benchmark stopping protocol. All observations are retained. The four architectures and their generated lengths differ; these data do not establish constant overhead or optimized-serving performance.

Model	Faithfulness		Factuality		Average
	RACE	SQuAD	NQ	TriviaQA	
Llama-3.2-3B	0.10	0.12	0.14	0.19	0.14
Llama-3.2-3B-Instruct	0.09	0.14	0.12	0.16	0.13
Llama-3-8B	0.09	0.12	0.14	0.16	0.13
Llama-3-8B-Instruct	0.09	0.12	0.16	0.16	0.13
Gemma-2-9B	0.11	0.10	0.12	0.12	0.11
Gemma-2-9B-Instruct	0.08	0.10	0.09	0.10	0.09
Average	0.09	0.12	0.13	0.15	0.12

Table 7: Threshold values for various settings using exact match as the correctness measure.

a threshold τ converts the score into a binary decision and summarize its observed variation in the original experiments.

Threshold Selection via Balanced Operating Point. Selecting an appropriate threshold τ involves balancing two competing objectives: sensitivity (detecting hallucinated generations) and specificity (avoiding false alarms on faithful outputs). Optimizing for accuracy alone can be misleading, particularly under class imbalance, as it may bias the threshold toward the majority class. To avoid this, we select τ by maximizing the Geometric Mean (G-Mean) of the True Positive Rate (TPR) and the True Negative Rate (TNR):

$$\text{G-Mean} = \sqrt{\text{TPR} \times \text{TNR}} = \sqrt{\text{TPR} \times (1 - \text{FPR})}.$$

We sweep τ over the range of observed HIDE scores on a validation set and select the value that maximizes G-Mean. This criterion penalizes extreme failure in either

direction and yields a balanced operating point, even when hallucinations constitute a minority of the data.

Variance of the Optimal Threshold. Beyond the consistency of the average optimal threshold ($\tau_{\text{avg}} \approx 0.12$), we further analyze the variance of τ across datasets and model variants to assess deployment stability. Overall, we observe that the optimal threshold remains largely consistent between base and instruction-tuned models. For example, Llama-3-8B exhibits an average optimal threshold of $\tau = 0.13$ for both its Base and Instruct variants, while Llama-3.2-3B shows similarly close values ($\tau = 0.14$ for Base and $\tau = 0.13$ for Instruct). Gemma-2-9B follows the same pattern, with $\tau = 0.11$ for the Base model and $\tau = 0.09$ for the Instruct variant. We also observe that across all four evaluated datasets, i.e., SQuAD, RACE, Natural Questions, and TriviaQA, the optimal thresholds are mostly consistent, as shown in Table 7, yielding an average threshold value $\tau_{\text{avg}} = 0.12$. These are similarities in the benchmark operating points; they do not establish a structural mechanism or transfer to the corrected protocol. We provide examples of a few input prompts from various datasets and the corresponding generations from Llama-3-8B along with the respective HIDE scores in Appendix D, to illustrate how hallucinated generations are detected by HIDE.

Analyzing variance across datasets reveals a more nuanced picture. Instruction-tuned models often exhibit lower threshold variance than their base counterparts – e.g., Llama-3.2-3B-Instruct shows a variance of 0.00087 compared to 0.00141 for its Base version, and Gemma-2-9B-Instruct shows a variance of 0.00005 versus 0.00012 for Gemma-2-9B Base. However, this trend is not universal: Llama-3-8B-Instruct exhibits slightly higher variance (0.00114) than its Base version (0.00087). These results indicate that while instruction tuning often stabilizes the optimal threshold, the effect is model-family dependent.

Taken together, these benchmark findings describe threshold variation within the evaluated protocol. Deployment under different generation, token-selection, or stopping conditions requires fresh validation-based calibration; the table does not establish a universally transferable threshold.

Interpreting τ in Deployment. The threshold τ selected during validation defines a default operating point for inference: generations with HIDE scores below τ are flagged as hallucinated. In practice, however, τ can be adjusted to reflect application-specific risk preferences. In safety-critical domains, such as medical or legal assistance, practitioners may raise τ to prioritize recall, accepting a higher false positive rate in order to minimize undetected hallucinations. Conversely, in user-experience-sensitive applications, such as creative writing or casual dialogue, a lower τ may be preferred to reduce unnecessary flags on correct outputs. Accordingly, while the G-Mean-optimized threshold provides a principled and balanced baseline, τ should be viewed as a tunable operating parameter rather than a fixed constant.

7. Analyses

To assess the robustness and practical implications of HIDE, we conduct a series of ablation studies that examine the impact of core methodological choices. These analyses investigate the following: (i) effectiveness of linear similarity measure, like cosine similarity, in capturing decoupling, (ii) impact of the number of selected tokens, (iii) dependence on the choice of transformer layer, (iv) the influence of kernel type and kernel hyperparameters, (v) the effect of biased versus our adapted HSIC estimators, and (vi)

Model	Method	SQuAD		NQ	
		AUC_s	PCC_s	AUC_s	PCC_s
Llama-3-8B	HSIC (HIDE)	76.87	49.09	79.02	49.61
	Cosine Similarity	49.78	8.78	65.54	25.43
Gemma-2-9B	HSIC (HIDE)	75.17	47.63	79.08	24.08
	Cosine Similarity	54.87	14.67	88.44	17.85

Table 8: Comparison of HSIC with cosine similarity for measuring input-output representational dependence on SQuAD and Natural Questions (NQ). HSIC consistently yields stronger ranking performance (AUC_s) and higher score-label correlation (PCC_s) on SQuAD, while cosine similarity exhibits unstable behavior across datasets and models.




Figure 5: Benchmark token-budget ablation on SQuAD and NQ using Llama-3-8B and Gemma-2-9B. The horizontal axis gives the configured selected-token budget; the vertical axis is AUC_s (sentence-similarity correctness threshold 0.9, higher is better). The actual selected count may be smaller for short answers. Curves level off around 15–20 tokens.

the relative merits of different token selection strategies. The results, summarized below, reinforce the robustness of HIDE across different design choices.

7.1 Non-linear Dependence vs. Linear Similarity

Table 8 compares HSIC with cosine similarity using both ranking-based (AUC_s) and correlation-based (PCC_s) metrics. On SQuAD, cosine similarity performs near random guessing, achieving low AUC_s and negligible correlation with ground-truth labels, whereas HSIC attains substantially higher values on both metrics across models. On Natural Questions, cosine similarity exhibits highly variable behavior across models — including strong AUC_s in isolated cases but weak correlation — while HSIC provides more stable performance across metrics. This benchmark comparison establishes performance relative to the evaluated cosine baseline, but does not establish that nonlinear dependence is necessary. The new attention and token-count controls qualify that interpretation (Section 6.2).




Figure 6: Benchmark layer ablation for Llama-3-8B and Gemma-2-9B on SQuAD and NQ. The horizontal axis gives the selected layer and the vertical axis AUC_s (sentence-similarity correctness threshold 0.9, higher is better). Curves show limited variation across the tested layers.

7.2 Impact of the Number of Selected Tokens

We begin by analyzing how the number of selected tokens (n_{eff}) affects detection performance. Figure 5 illustrates that as n_{eff} increases from 5 to 10, there is a clear improvement in the AUC_s scores, but further increases beyond 15 – 20 tokens result in diminishing returns. This levelling-off happened earlier for NQ across both models, likely because its shorter outputs often meant the actual number of tokens used was limited, regardless of a higher token budget. This finding holds consistently across datasets (SQuAD and NQ) and models (Llama-3-8B and Gemma-2-9B), indicating that HIDE achieves near-optimal performance with a modest sample size. This property is practically significant, as it enables computational efficiency without sacrificing detection quality, even for outputs of varying lengths.

7.3 Robustness to Selected Layer

We evaluate all 32 layers of Llama-3-8B and all 42 layers of Gemma-2-9B. Figure 6 shows little variation in AUC_s , with a mild optimum near the middle layers. This supports limited sensitivity to layer choice in these experiments. It does not show that the score detects the same semantic information at every layer; the selected-token-count result in Section 6.2 offers another explanation for the stability.

7.4 Effect of Kernel Function and RBF Hyperparameter

Since HIDE quantifies statistical dependence using HSIC, the choice of kernel function and its parameters is a critical consideration. We evaluate several kernel types – including RBF (default), linear, polynomial, cosine, sigmoid, laplacian, exponential, periodic, and matern – as well as a range of γ values for the RBF kernel.

Figure 7a demonstrates that, with the exception of the polynomial kernel (which underperforms substantially), most kernels yield robust and similar AUC_s scores for SQuAD and NQ, using Llama-3-8B. This suggests that the effectiveness of HIDE is largely preserved regardless of kernel choice, with all kernels, except polynomial, providing strong performance.

For the RBF kernel, Figure 7b shows stable performance for γ from 10^{-9} to 10^{-3} and a sharp decline for $\gamma \geq 10^{-1}$. Larger values make the kernel more sensitive to distances




Figure 7: Comparison of AUC_s scores across (a) different kernel functions, and (b) different values of the hyperparameter γ for the RBF kernel used for HIDE-score calculation on the SQuAD and NQ datasets using Llama-3-8B.




Figure 8: Benchmark estimator comparison for Llama-3-8B and Gemma-2-9B on SQuAD and NQ. Bars compare the adapted HIDE estimator with biased HSIC using AUC_s (sentence-similarity correctness threshold 0.9, higher is better). Each bar reports the corresponding model/dataset result under the benchmark generation protocol.

between hidden states. These results describe sensitivity to the kernel parameter, but do not establish why the default value predicts correctness. The selected-token-count comparison in Section 6.2 helps assess that question.

7.5 Effectiveness of HSIC Estimators

We compare the performance of HIDE with the biased estimator, $\widehat{HSIC}_b$, and our adapted estimator $\widehat{HSIC}_{HIDE}$, for empirical HSIC computation, as shown in Figure 8. $\widehat{HSIC}_{HIDE}$ delivers consistently higher performance across datasets and models, owing to improved score separability between hallucinated and faithful generations. As shown in Figure 9, $\widehat{HSIC}_{HIDE}$ produces a more favourable score distribution spanning 0 – 0.25, creating clearer separation between hallucinated and non-hallucinated generations. In contrast, $\widehat{HSIC}_b$ yields extremely small values, compressing scores close to zero (with a maximum score of 2.74×10^{-11}), making threshold determination challenging and less reliable. Therefore, we adopt $\widehat{HSIC}_{HIDE}$ for our framework.




Figure 9: Distribution of HIDE scores using $\widehat{\text{HSIC}}_b$ and $\widehat{\text{HSIC}}_{\text{HIDE}}$ estimators for HSIC calculation, for Llama-3-8B on SQuAD and NQ datasets.




Figure 10: Comparison of AUC_s scores for keyword-based and SVD-based token selection strategies, for SQuAD and NQ datasets using Llama-3-8B and Gemma-2-9B.

7.6 Impact of Token Selection Strategies

Finally, we compared two strategies for selecting token representations (as discussed in §4.3.1): keyword-based selection and SVD-based alignment. Figure 10 shows that both approaches perform nearly identically across datasets and architectures. The SVD method guarantees mathematical orthogonality, while the keyword-based approach enhances interpretability by enabling semantic traceability – allowing users to attribute detection decisions to specific content elements. Given this enhanced transparency of the keyword-based strategy and the absence of any performance penalty, we adopt the keyword-based selection as the default for HIDE.

7.7 Qualitative Analysis of Representational Decoupling

Table 9 illustrates how representational decoupling aligns with hallucination behavior. Faithful generations preserve alignment between salient input and output concepts, yielding high HIDE scores. Hallucinated outputs introduce unsupported entities or semantic drift, resulting in substantially lower dependence scores. This contrast provides qualitative intuition for why HSIC-based dependence captures hallucinations beyond surface-level similarity. Several more illustrative examples from different datasets are presented in Appendix D.

Prompt (Dataset)	Input words	Key-words	Output Keywords	Model Output	Gold swer	An- HIDE Score	Hallucinated?
Who had a 70s hit with Kiss? You All Over? (TriviaQA)	70s No.1 Kiss song	hit, Kiss	Exile, band, hit	Exile	Exile	0.182	No
What claimed the life of singer Kathleen Ferrier? (TriviaQA)	Kathleen Ferrier, singer, death		cancer, illness	Cancer	Cancer	0.164	No
Who was the man behind The Chipmunks? (TriviaQA)	The, Chipmunks, Who, ...		Bagda-sarian, Songbook, Ross, name, first, Sr, ...	Ross Bagda-sarian, Sr.	David Seville	0.066	Yes
What is the capital of Australia? (NQ)	Australia, of capital		Sydney, city	Sydney	Canberra	0.034	Yes

Table 9: Benchmark qualitative examples for Llama-3-8B, with input/output keywords, generated answers, first reference, and HIDE score. The displayed labels follow the original reference-based evaluation; aliases and surface-form differences can affect these labels. Selected examples do not establish a causal explanation of score differences.

7.8 Robustness to Decoding Strategies

Autoregressive internal dynamics and hallucination rates are known to vary across different decoding regimes. Theoretically, HIDE’s effectiveness is decoupled from the specific token-selection mechanism. If a stochastic sampling choice forces the generation off-topic or into a parametric hallucination, the output representations structurally diverge from the input context, and the HSIC score should reliably reflect this decoupling.

To empirically validate this, we evaluated HIDE across multiple temperature ($T \in \{0.3, 0.6, 0.9\}$) and nucleus sampling ($Top-p \in \{0.7, 0.8, 0.9\}$) configurations on SQuAD and Natural Questions. As detailed in Table 10, HIDE exhibits remarkable stability across all settings. The AUC_s and PCC_s scores remain consistently strong, and in several stochastic configurations, performance marginally improves over the greedy baseline. This confirms that representational decoupling is a robust, generation-agnostic signal.

8. Error Analysis

While HIDE demonstrates robust performance across diverse hallucination detection settings, our detailed error analysis reveals specific cases where the method faces limitations, providing guidance for future improvements. Here, we discuss the principal sources of error and critically assess their implications for practical deployment.

8.1 Insufficient Token Selection in Single Word-Answer Tasks

Model	Decoding Strategy	SQuAD (Faithfulness)		NQ (Factuality)	
		AUC _s	PCC _s	AUC _s	PCC _s
Llama-3-8B	Greedy	76.87	49.09	79.02	49.61
	$T = 0.3$	77.02	49.61	79.32	49.93
	$T = 0.6$	79.58	52.44	77.90	45.64
	$T = 0.9$	80.74	53.02	81.37	48.40
	$Top-p = 0.9$	79.63	51.17	78.12	44.68
	$Top-p = 0.8$	79.09	50.70	81.82	46.85
	$Top-p = 0.7$	79.62	52.35	80.13	47.79
Gemma-2-9B	Greedy	75.17	47.63	79.08	24.08
	$T = 0.3$	75.72	47.48	80.97	34.85
	$T = 0.6$	76.89	48.46	81.84	42.70
	$T = 0.9$	76.41	46.20	78.28	44.21
	$Top-p = 0.9$	77.95	47.22	75.88	42.76
	$Top-p = 0.8$	78.09	47.41	77.84	42.86
	$Top-p = 0.7$	77.15	47.79	76.41	44.04

Table 10: Robustness of HIDE across various decoding regimes. Performance remains highly stable across greedy and stochastic generation strategies.

Very short responses leave few output states available for scoring. When only one entry is selected from each sequence, the adapted score is zero, as shown in Lemma 3 of Appendix A.3. If no output state is available, the implementation also assigns zero. These cases are common concerns for short factual answers such as those in TriviaQA. For instance, given the prompt, “Answer these questions: Q: Captain Corelli’s mandolin is a book and film set in which country? A:” from TriviaQA, Llama-3-8B generates the output “Greece”, containing just a single token. Here, though the output is correct and matches the ground-truth, HIDE cannot meaningfully assess input-output dependence, resulting in an uninformative score of zero. This limitation highlights the need for additional heuristics or alternative strategies when handling extremely small outputs, as the statistical underpinning of HIDE makes it most effective with moderately long sequences. Importantly, this behavior is not an artefact of our implementation but a consequence of the underlying statistics. HIDE measures statistical dependence between input and output representations, which is a distributional property requiring multiple samples to estimate variance and covariance. As the output length approaches a single token, dependence estimation becomes mathematically ill-defined. In this regime, uncertainty-based measures such as perplexity provide the appropriate signal, as they directly quantify the model’s confidence in a single prediction. Thus, reverting to perplexity for $|S_{\text{out}}| = 1$ is not an ad-hoc patch, but a principled transition between statistical regimes. A possible way to handle this in real-world deployment is therefore to fall back to uncertainty-based single-pass detection methods, such as perplexity, in the case of single-token outputs. As also noted by Chen et al. (2024a), simple uncertainty-based measures like perplexity work well for datasets like TriviaQA where the expected output is very short, typically consisting of one or two tokens. We quantify the prevalence of such very short outputs across QA benchmarks and analyze their practical impact in Appendix E.

8.2 Degraded Performance Due to Input Prompt Copying

HIDE also exhibits weakness in scenarios where the language model’s outputs contain verbatim repetitions of the input prompt. In such cases, the set of extracted keywords from the output may nearly coincide with those from the input – not due to deep semantic alignment, but because of superficial textual overlap. For instance, for a prompt “*Who was the target of the failed “Bomb Plot” of 1944?*” from TriviaQA, the generated output from Llama-3-8B goes as “*The target of the failed “Bomb Plot” of 1944 was Adolf Hitler. Q: Who was the target of the failed “Bomb Plot” of 1944? ...*”, which contains verbatim **repetition** of the input question along with the answer, thereby artificially inflating the computed statistical dependence. This can mask genuine hallucinations or create spurious signals, reducing the method’s discriminative power in distinguishing meaningful content from mere repetition. Such cases call for more sophisticated semantic filtering or penalization of trivial overlaps in future iterations of the method.

9. Conclusions

We introduced HIDE, a training-free detector computed from selected input and output hidden states during one generation. Across six models and four QA datasets, it is competitive with the evaluated baselines. It reduces mean latency by 50.8% relative to five-generation EigenScore under standard HuggingFace inference. Attention performs better in seven of eight directly compared settings; HIDE performs better on Gemma-2-9B/RACE. The selected-token-count control closely tracks HIDE, limiting a nonlinear-dependence explanation of its performance. Timing measurements up to Gemma-2-27B show that the added overhead varies with model and dataset. These results show where HIDE is useful and where simpler signals perform better.

Limitations

The adapted score depends strongly on the number of selected tokens. For all-ones input and output Gram matrices, it equals $(n - 1)/n^2$ for $n \geq 1$, rather than zero in general. Its agreement with this control limits a semantic-dependence interpretation. The assumptions of the asymptotic result in Appendix A.3 need not hold for selected tokens within a single answer.

The full-benchmark tables use stopping conditions that can allow generation to continue beyond an answer. This can affect correctness labels and scores. The first-answer-line comparison avoids this source of variation, but covers only two models and 2,000 examples per dataset. Its sample size was chosen after earlier runs had been inspected, although examples were not selected by outcome. It does not validate the full-benchmark results under these stopping conditions. Sentence similarity and ROUGE-L are imperfect measures of correctness, and the evaluation uses only the first reference answer.

HIDE requires access to hidden states; the attention baseline requires attention weights. The norm baseline omits value and output projections and is therefore only a proxy for the input-derived update. Correlations and the QCBW comparison cannot identify causal mechanisms. The timing experiments use standard inference and benchmark stopping conditions; they do not measure optimized serving or the first-answer-line protocol. Host sharing may affect the measurements, and the range of hidden widths is limited even with Gemma-2-27B.

Appendix A: Theoretical Foundations of HSIC and HIDE

A.1 Measuring Statistical Dependence using HSIC

Let's assume that we want to assess the statistical dependence between two random variables X and Y , taking values in separable topological spaces $\mathcal{X}$ and $\mathcal{Y}$, respectively. The statistical relationship between X and Y is governed by their joint probability distribution P_{XY} . HSIC quantifies this dependence by comparing P_{XY} to the product of the marginals, $P_X P_Y$, within Reproducing Kernel Hilbert Spaces (RKHSs). An RKHS $\mathcal{F}_k$, associated with a kernel $k : \mathcal{X} \times \mathcal{X} \rightarrow \mathbb{R}$ allows implicit mapping to a feature space via $\phi : \mathcal{X} \rightarrow \mathcal{F}_k$, where $k(x, x') = \langle \phi(x), \phi(x') \rangle_{\mathcal{F}_k}$, with $\langle \cdot, \cdot \rangle_{\mathcal{F}_k}$ denoting the inner product in the RKHS $\mathcal{F}_k$. This mapping facilitates the detection of non-linear dependencies.

Intuitively, standard linear metrics (like cosine similarity) only capture direct geometric alignment. In contrast, evaluating dependence in an RKHS allows us to detect arbitrary, non-linear relationships. The cross-covariance operator, C_{XY} , acts as a high-dimensional analog to the standard covariance matrix, capturing the covariance between the transformed representations of X and Y in their respective feature spaces. A reduced cross-covariance magnitude motivates the decoupling hypothesis, but the observed score and proxy correlations do not establish this as a general consequence or cause of hallucination. Formally, given two RKHSs, $\mathcal{F}$ with kernel k_X for variable X , and $\mathcal{G}$ with kernel k_Y for variable Y (corresponding to feature maps ϕ and ψ respectively), the cross-covariance operator $C_{XY} : \mathcal{G} \rightarrow \mathcal{F}$ is defined as:

$$C_{XY} := \mathbb{E}_{XY}[\phi(X) \otimes \psi(Y)] - \mathbb{E}_X[\phi(X)] \otimes \mathbb{E}_Y[\psi(Y)]$$

where $\otimes$ denotes the tensor product. If X and Y are independent, C_{XY} is the zero operator.

The Hilbert-Schmidt norm of an operator $C : \mathcal{G} \rightarrow \mathcal{F}$, denoted by $\|C\|_{HS}$, is a generalization of the Frobenius norm for matrices and can be defined as $\|C\|_{HS}^2 = \sum_{i,j} |\langle Cg_j, f_i \rangle_{\mathcal{F}}|^2$, where $\{f_i\}$ and $\{g_j\}$ are orthonormal bases of $\mathcal{F}$ and $\mathcal{G}$, respectively. It essentially measures the *magnitude* of the operator.

Definition 3

HSIC is the squared Hilbert-Schmidt norm of the cross-covariance operator C_{XY} :

$$HSIC(X, Y; k_X, k_Y) := \|C_{XY}\|_{HS}^2 \quad (\text{A.1})$$

A crucial property of HSIC is its relationship with statistical independence, formalized as follows:

Lemma 1 (Gretton et al. (2005))

Let k_X and k_Y be characteristic kernels on $\mathcal{X}$ and $\mathcal{Y}$, respectively. Then $HSIC(X, Y; k_X, k_Y) = 0$ if and only if X and Y are statistically independent.

A kernel k is termed characteristic if the mapping from probability measures on its domain to mean elements in its RKHS, $\mu \mapsto \mathbb{E}_{X \sim \mu}[\phi(X)]$, is injective (Sriperumbudur et al. 2010). This ensures that the mean embedding $\mathbb{E}[\phi(X)]$ uniquely determines P_X . For our method, we use the Radial Basis Function (RBF) kernel, $k_{\text{RBF}}(x, x') = \exp(-\gamma \|x - x'\|_2^2)$ for a bandwidth parameter $\gamma > 0$.

Lemma 2 (Sriperumbudur et al. (2008))

The RBF kernel is characteristic on $\mathbb{R}^d$.

The conjunction of Lemma 1 and Lemma 2 establishes HSIC, when computed with RBF kernels, as a robust and theoretically sound metric for detecting statistical independence, forming the cornerstone of our approach.

A.2 Empirical Estimation of HSIC

In practice, we rely on empirical estimates of HSIC to quantify statistical independence from a finite set of n samples. Given samples $\{x_1, \dots, x_n\}$ from the distribution of X and $\{y_1, \dots, y_n\}$ from the distribution of Y , we form $n \times n$ Gram matrices $\mathbf{K}_X$ and $\mathbf{K}_Y$, where $\mathbf{K}_{X,ij} = k_X(x_i, x_j)$ and $\mathbf{K}_{Y,ij} = k_Y(y_i, y_j)$.

A common, though biased, empirical estimator for HSIC is (Gretton et al. 2005):

$$\widehat{\text{HSIC}}_b(X, Y) = \frac{1}{(n-1)^2} \text{Tr}(\mathbf{K}_X \mathbf{H} \mathbf{K}_Y \mathbf{H}) \quad (\text{A.2})$$

where $\mathbf{H} = \mathbf{I} - \frac{1}{n} \mathbf{1} \mathbf{1}^T$ is the $n \times n$ centering matrix, $\mathbf{I}$ is the identity matrix, and $\mathbf{1}$ is an $n \times 1$ column vector of ones.

For improved performance, particularly with smaller n , an unbiased estimator is preferred (Song et al. 2012). **The unbiased estimator is given by:**

$$\widehat{\text{HSIC}}_u(X, Y) = \frac{1}{n(n-3)} \left[\text{Tr}(\tilde{\mathbf{K}}_X \tilde{\mathbf{K}}_Y) + \frac{(\mathbf{1}^T \tilde{\mathbf{K}}_X \mathbf{1})(\mathbf{1}^T \tilde{\mathbf{K}}_Y \mathbf{1})}{(n-1)(n-2)} - \frac{2}{n-2} \mathbf{1}^T \tilde{\mathbf{K}}_X \tilde{\mathbf{K}}_Y \mathbf{1} \right] \quad (\text{A.3})$$

Here, n is the number of samples. The matrices $\tilde{\mathbf{K}}_X$ and $\tilde{\mathbf{K}}_Y$ are derived from $\mathbf{K}_X$ and $\mathbf{K}_Y$, respectively by setting their diagonal entries to zero: $\tilde{\mathbf{K}}_{X,ij} = (1 - \delta_{ij}) \mathbf{K}_{X,ij}$, where δ_{ij} is the Kronecker delta (1 if $i = j$, and 0 otherwise). This estimator requires $n \geq 4$ for its standard definition.

A.3 Properties of the HIDE Score

We now analyze the properties of our proposed HIDE score (Equation 1), denoted as $\widehat{\text{HSIC}}_{\text{HIDE}}$, **which adapts the standard unbiased HSIC estimator** to be computable for any number of effective tokens $n \equiv n_{\text{eff}} \geq 1$.

Let $\{(x_i, y_i)\}_{i=1}^n$ be the set of hidden state representations for the selected input and output tokens. We use the RBF kernel, which is bounded such that $0 < k(u, v) \leq 1$ for any vectors u, v . The score is computed using Gram matrices with zeroed diagonals, $\tilde{\mathbf{K}}_X$ and $\tilde{\mathbf{K}}_Y$, as defined in Equation A.3.

Lemma 3 (Behaviour for Small Sample Sizes)

The $\widehat{\text{HSIC}}_{\text{HIDE}}$ estimator is computable for small sample sizes where $\widehat{\text{HSIC}}_u$ is undefined. Specifically:

- (i) If $n = 1$, then $\widehat{\text{HSIC}}_{\text{HIDE}}(X, Y) = 0$.
- (ii) If $n = 2$, then $\widehat{\text{HSIC}}_{\text{HIDE}}(X, Y) = \frac{1}{4} k_X(x_1, x_2) k_Y(y_1, y_2)$, which implies $0 < \widehat{\text{HSIC}}_{\text{HIDE}}(X, Y) \leq 1/4$.

Proof. The proof of the above lemma follows from direct substitution into Equation 1 for computing $\widehat{\text{HSIC}}_{\text{HIDE}}$.

(i) For $n = 1$, the zero-diagonal matrices are $\tilde{K}_X = [0]$ and $\tilde{K}_Y = [0]$. All terms in Equation 1 thus becomes zero, implying $\widehat{\text{HSIC}}_{\text{HIDE}} = 0$.

(ii) For $n = 2$, let $a = k_X(x_1, x_2)$ and $b = k_Y(y_1, y_2)$. The matrices are, therefore, $\tilde{K}_X = \begin{pmatrix} 0 & a \\ a & 0 \end{pmatrix}$ and $\tilde{K}_Y = \begin{pmatrix} 0 & b \\ b & 0 \end{pmatrix}$. The constituent terms of Equation 1 are: $\text{Tr}(\tilde{K}_X \tilde{K}_Y) = 2ab$; $\mathbf{1}^\top \tilde{K}_X \mathbf{1} = 2a$; $\mathbf{1}^\top \tilde{K}_Y \mathbf{1} = 2b$; and $\mathbf{1}^\top \tilde{K}_X \tilde{K}_Y \mathbf{1} = 2ab$. Substituting these into Equation 1 for $n = 2$ yields:

$$\widehat{\text{HSIC}}_{\text{HIDE}}(X, Y) = \frac{1}{2^2} \left[2ab + \frac{(2a)(2b)}{2^2} - \frac{2}{2}(2ab) \right] = \frac{1}{4} [2ab + ab - 2ab] = \frac{ab}{4}.$$

Since the RBF kernel ensures $a, b \in (0, 1]$ for distinct points, the result $(0, 1/4]$ follows. This confirms the score yields a deterministic, non-negative value for $n = 2$. $\square$

Lemma 4 (Asymptotic Properties)

For independent and identically distributed paired observations (x_i, y_i) and fixed bounded RBF kernels, the $\widehat{\text{HSIC}}_{\text{HIDE}}$ estimator is asymptotically unbiased and strongly consistent.

(i) (Asymptotic Bias) *The bias of the estimator vanishes as $n \rightarrow \infty$:*

$$|\mathbb{E}[\widehat{\text{HSIC}}_{\text{HIDE}}] - \text{HSIC}(X, Y)| = O(1/n).$$

(ii) (Strong Consistency) *As $n \rightarrow \infty$, $\widehat{\text{HSIC}}_{\text{HIDE}}(X, Y)$ converges almost surely to the true population HSIC:*

$$\widehat{\text{HSIC}}_{\text{HIDE}}(X, Y) \xrightarrow{a.s.} \text{HSIC}(X, Y).$$

Proof. Write $A_n = \text{Tr}(\tilde{K}_X \tilde{K}_Y)$, $B_n = (\mathbf{1}^\top \tilde{K}_X \mathbf{1})(\mathbf{1}^\top \tilde{K}_Y \mathbf{1})$, and $C_n = \mathbf{1}^\top \tilde{K}_X \tilde{K}_Y \mathbf{1}$. Since the kernel entries are bounded by one, $|A_n| \leq n^2$, $|B_n| \leq n^4$, and $|C_n| \leq n^3$. The differences between the coefficients of these three terms in the adapted and unbiased estimators satisfy, as $n \rightarrow \infty$,

$$\begin{aligned} \left| \frac{1}{n^2} - \frac{1}{n(n-3)} \right| &= O(n^{-3}), \\ \left| \frac{1}{n^4} - \frac{1}{n(n-3)(n-1)(n-2)} \right| &= O(n^{-5}), \\ \left| \frac{2}{n^3} - \frac{2}{n(n-3)(n-2)} \right| &= O(n^{-4}). \end{aligned}$$

Multiplying these bounds gives $|\widehat{\text{HSIC}}_{\text{HIDE}} - \widehat{\text{HSIC}}_u| = O(1/n)$, uniformly over the observations. Since the standard estimator is unbiased under the stated sampling assumptions, taking expectations proves (i). Its strong consistency under the same assumptions, together with the vanishing uniform difference, proves (ii). This limiting result does not imply unbiasedness at a fixed small token count, nor does it establish independence for token representations selected from a single generated answer. $\square$

Appendix B: Detailed Per-Dataset Results

For completeness, this appendix reports the full per-dataset hallucination detection results underlying the summarized comparisons in the main paper. Table B.1 presents AUC-ROC and PCC scores for all methods across individual datasets and models, allowing fine-grained inspection of dataset-specific behavior and variance that is abstracted away in the averaged results shown in Table 2.

Appendix C: Performance with Exact Match as Correctness Measure

In addition to AUC-ROC and PCC for sentence embedding similarity and ROUGE-L as correctness measures, we also present the evaluation results by using exact match (Liang et al. 2023) as the correctness measure, which imposes a more stringent criterion for a generation to be accurate. The findings in Table C.1 provide results analogous to those in Table B.1, indicating that our approach is effective for detecting both faithfulness and factuality hallucinations.

Appendix D: Illustrative Examples

We provide some illustrative examples in this section for prompts from various datasets and the corresponding outputs from Llama-3-8B with the HIDE scores for every generation. We observe a threshold of 0.12, as discussed in §6.5, effectively discriminates hallucinated generations in most cases.

Example from SQuAD

Prompt: Civil disobedience is usually defined as pertaining to a citizen’s relation to the state and its laws, as distinguished from a constitutional impasse in which two public agencies, especially two equally sovereign branches of government, conflict. For instance, if the head of government of a country were to refuse to enforce a decision of that country’s highest court, it would not be civil disobedience, since the head of government would be acting in her or his capacity as public official rather than private citizen.

Q: What type of person can not be attributed civil disobedience?

A:

Ground-truth: head of government

Llama-3-8B Output: A person who is not a citizen of the country in which the civil disobedience is taking place.

Input Keywords: [disobedience, civil, person, enforce, attributed, defined, agencies, relation, be, head, his, official, usually, instance, capacity, impasse, would, branches, highest, and]

Output Keywords: [disobedience, citizen, civil, person, country, who, A, a, the, place, which, taking, in, is, of, not]

Selected Input Tokens: [disobed, ience, disobed, ience, disobed, ience, civil, civil, person, enforce, attributed, defined, agencies, relation, be, be, be, head, head]

Selected Output Tokens: [disobed, ience, citizen, civil, person, country, who, A, a, the, the, place, which, taking, in, is, is, of, not]

HIDE Score: 0.049

Table B.1: Full-benchmark AUC-ROC and PCC values under the original stopping protocol, using sentence similarity and ROUGE-L as correctness measures for faithfulness and factuality hallucination detection across four datasets. HIDE is compared with single-pass and multi-pass baselines (rows corresponding to multi-pass are highlighted in gray; five outputs are generated for each multi-pass method). The best result for each metric is shown in **bold**; the best among single-pass methods is underlined.

Method	Faithfulness						Factuality						Average							
	RACE			SQuAD			NQ			TriviaQA			Average							
	AUC _c	AUC _f	PCC _c	PCC _f	AUC _c	AUC _f	PCC _c	PCC _f	AUC _c	AUC _f	PCC _c	PCC _f	AUC _c	AUC _f	PCC _c	PCC _f				
Perplexity Energy LN-Entropy Lexical Similarity Eigenscore HIDE	50.27	49.94	18.6	15	46.53	47.3	-7.81	-1.56	53.03	54.25	-8.16	9.2	57.53	57.01	6.14	20.41	51.84	52.13	2.19	10.76
	49.07	49.93	13.10	9.14	36.9	37.22	-29.49	-26.22	36.83	35.2	-25.81	-10.75	43.32	41.84	-34.53	-22.20	41.53	41.05	-19.18	-12.51
	50.90	50.26	6.01	6.96	60.85	60.89	14.24	19.74	72.39	71.31	20.58	22.53	67.31	67.24	22.52	28.82	62.86	62.42	15.84	19.52
	64.11	65.04	22.04	27.32	64.61	66.71	28.81	34.1	65.98	69.51	20.32	27.8	79.74	81.34	52.99	56.59	68.61	70.65	31.04	36.45
	58.17	57.18	9.56	12.37	65.31	65.49	21.29	25.46	68.34	67.34	20.27	17.39	76.13	76.57	50.45	48.15	66.65	65.39	25.85	25.85
	66.43	65.48	18.34	20.72	77.55	77.28	49.86	50.3	75.16	74.63	47.77	25.27	58.36	59.14	27.46	17.88	69.37	69.13	35.86	28.67
Perplexity Energy LN-Entropy Lexical Similarity Eigenscore HIDE	59.52	61.28	34.59	28.44	52.50	55.22	9.73	15.36	54.90	54.16	-7.26	8.80	60.08	60.14	6.42	22.49	56.75	57.70	10.87	18.77
	54.19	57.03	21.48	17.04	37.36	38.57	-18.23	-16.81	25.89	24.72	-48.16	-25.99	42.51	41.59	-31.54	-16.17	39.99	40.48	-19.11	-10.48
	54.46	55.65	14.96	12.96	62.98	64.74	23.51	29.69	71.88	70.42	25.68	24.93	71.83	72.01	22.95	29.18	65.70	65.70	21.77	24.19
	66.75	68.20	31.91	33.83	69.05	71.69	35.08	41.91	70.49	70.53	34.64	32.62	79.25	80.59	51.98	54.42	71.38	72.75	38.38	40.69
	62.82	62.09	19.09	20.92	73.14	73.57	39.38	45.56	76.42	75.56	41.64	32.71	78.35	79.08	53.46	49.55	72.68	72.57	38.44	37.19
	74.88	73.02	27.36	33.95	75.18	75.48	42.80	44.98	76.56	77.27	53.91	35.40	58.43	59.03	50.64	19.54	71.26	71.20	38.68	33.47
Perplexity Energy LN-Entropy Lexical Similarity Eigenscore HIDE	50.78	51.93	23.01	18.07	45.08	46.37	-8.01	-2.39	48.18	47.51	-11.22	2.16	38.21	36.90	-19.49	-8.54	45.56	45.68	-3.92	2.33
	50.10	51.35	17.68	12.31	35.27	35.48	-29.16	-26.95	33.70	32.13	-23.86	-15.02	27.43	26.48	-50.84	-42.86	36.63	36.36	-21.55	-18.13
	45.49	45.85	-0.10	-0.54	60.50	60.57	15.09	19.82	65.81	65.00	13.80	16.17	72.15	72.08	33.54	37.79	60.99	60.87	15.58	18.31
	66.10	66.92	29.15	33.80	64.01	66.57	29.36	33.69	65.22	65.84	20.10	25.87	77.13	78.65	47.18	50.24	68.12	69.50	31.45	35.90
	54.16	53.98	6.73	10.88	66.54	66.34	25.66	29.93	64.36	64.76	17.07	16.90	81.06	82.14	57.92	57.46	66.81	66.84	26.84	28.79
	69.81	66.95	16.76	22.87	76.87	77.45	49.09	50.43	79.02	79.17	49.61	33.91	65.65	66.82	41.82	35.19	72.84	72.60	39.32	35.60
Perplexity Energy LN-Entropy Lexical Similarity Eigenscore HIDE	60.20	61.91	22.52	22.42	47.46	50.73	2.95	6.94	57.47	56.11	-7.64	9.72	61.70	61.14	7.68	21.43	56.71	57.47	6.38	15.13
	55.94	57.16	18.26	15.97	43.05	44.60	-8.96	-8.31	30.63	29.56	-36.02	-18.50	39.40	38.80	-29.22	-16.69	42.25	42.53	-13.99	-6.88
	57.13	57.83	12.59	14.26	62.73	63.44	23.11	26.94	67.29	65.43	18.12	20.80	79.61	79.72	42.76	48.41	66.69	66.60	24.15	27.60
	69.28	71.22	36.64	40.07	62.46	65.14	25.76	29.69	69.05	68.88	26.24	28.84	81.12	82.33	58.68	59.19	70.48	71.89	36.83	39.45
	63.65	62.99	18.78	23.01	72.11	71.49	36.95	39.85	77.89	77.53	41.23	32.73	82.83	83.94	69.32	63.72	74.12	73.99	41.57	39.83
	73.20	72.73	22.12	30.11	80.15	79.14	52.27	55.98	79.29	78.91	56.33	37.98	61.70	62.94	36.73	27.27	73.58	73.43	41.86	37.84
Perplexity Energy LN-Entropy Lexical Similarity Eigenscore HIDE	59.88	60.65	29.89	27.98	47.42	48.29	-8.61	-1.85	29.92	35.59	9.29	12.95	43.08	46.18	-6.69	6.65	45.07	47.68	5.97	11.43
	44.05	43.20	-13.38	-12.32	52.12	52.50	6.81	2.25	64.89	60.48	0.50	-3.93	68.90	67.80	14.82	7.61	57.49	56.00	2.19	-1.60
	63.27	64.11	25.93	26.76	60.05	61.26	15.72	22.29	44.14	53.01	8.19	20.20	61.40	61.45	44.85	0.95	13.07	52.25	55.81	12.70
	69.02	69.62	34.18	35.48	67.10	70.32	35.25	40.56	56.80	59.43	13.50	23.50	62.98	65.46	23.79	32.99	63.97	66.21	26.68	33.13
	64.58	64.35	24.99	25.20	73.40	75.14	41.66	47.29	73.31	74.31	7.70	16.79	71.40	71.70	23.15	27.42	70.67	71.37	24.37	29.18
	58.06	57.99	20.60	16.50	75.17	76.66	47.63	48.43	79.08	75.96	24.08	15.41	80.84	81.02	50.51	43.88	73.29	72.91	35.70	31.05
Perplexity Energy LN-Entropy Lexical Similarity Eigenscore HIDE	60.22	61.27	14.11	17.12	53.15	56.90	5.73	12.49	31.52	34.14	29.54	0.19	32.09	36.85	52.60	11.92	44.25	47.29	25.49	10.43
	57.29	55.87	7.28	5.99	80.08	77.41	49.54	50.08	77.19	76.56	-9.51	11.55	92.34	90.96	-12.88	28.05	76.73	75.20	8.61	23.92
	62.82	64.58	17.67	20.84	67.09	69.50	26.65	33.47	44.68	46.01	13.79	4.64	39.70	44.13	28.49	14.62	53.57	56.06	21.65	18.39
	72.55	72.98	34.16	38.63	72.33	74.97	39.77	46.57	53.27	56.91	19.80	11.13	51.56	56.22	31.36	21.47	62.43	65.27	21.27	29.45
	70.22	70.59	24.78	29.44	71.70	73.99	44.71	49.87	65.97	67.11	8.40	14.46	61.17	64.63	15.05	21.13	67.26	69.08	23.24	28.73
	64.18	64.03	11.92	16.26	72.69	73.32	39.73	39.70	85.19	85.89	51.20	41.99	70.90	70.72	52.64	45.90	73.24	73.49	38.87	35.97

Method	Faithfulness				Factuality				Average	
	RACE		SQuAD		NQ		TriviaQA			
	AUC	PCC	AUC	PCC	AUC	PCC	AUC	PCC	AUC	PCC
Llama-3.2-3B										
Perplexity	52.91	7.32	45.32	-4.34	54.75	5.37	57.84	19.44	52.70	6.95
Energy	49.86	0.05	34.40	-23.90	35.01	-13.17	42.54	-18.99	40.45	-14.00
LN-Entropy	50.78	2.55	61.76	17.66	74.54	21.42	67.24	27.11	63.58	17.18
Lexical Similarity	66.17	24.29	64.13	22.29	68.86	22.29	79.32	50.78	69.62	29.91
Eigenscore	62.00	16.80	65.91	21.57	70.12	19.08	75.95	43.92	68.50	25.34
HIDE	69.63	26.13	77.30	44.15	74.46	26.23	57.61	14.08	69.75	27.65
Llama-3.2-3B-Instruct										
Perplexity	58.68	14.88	52.55	6.22	55.66	5.48	59.90	20.39	56.70	11.74
Energy	51.55	3.08	36.13	-21.11	25.59	-23.91	42.06	-14.11	38.83	-14.01
LN-Entropy	53.50	6.15	64.53	23.31	71.06	20.88	71.77	27.32	65.22	19.41
Lexical Similarity	66.56	27.05	69.52	31.73	70.58	27.62	78.88	49.17	71.38	33.89
Eigenscore	64.16	22.86	75.66	42.81	76.05	30.44	77.91	45.54	73.44	35.41
HIDE	78.49	41.68	75.71	42.45	75.55	32.40	57.23	14.99	71.74	32.88
Llama-3-8B										
Perplexity	56.68	11.82	44.16	-6.38	46.96	-1.38	39.23	-5.94	46.76	-0.47
Energy	55.98	9.88	33.51	-26.30	30.93	-18.07	27.31	-39.28	36.93	-18.44
LN-Entropy	45.33	-3.66	61.24	17.13	66.40	15.30	72.40	36.05	61.34	16.21
Lexical Similarity	66.05	25.92	63.41	22.28	65.05	20.06	76.29	44.85	67.70	28.28
Eigenscore	55.49	11.51	68.38	28.24	66.40	17.54	80.04	52.15	67.58	27.36
HIDE	69.48	24.27	76.57	44.64	78.58	33.65	64.35	28.95	72.24	32.88
Llama-3-8B-Instruct										
Perplexity	58.16	10.67	46.38	-2.87	57.74	6.97	60.95	18.78	55.81	8.39
Energy	52.26	3.91	41.37	-14.63	30.19	-19.62	38.32	-17.71	40.53	-12.01
LN-Entropy	57.33	11.06	63.65	19.21	67.65	19.26	78.61	43.82	66.81	23.34
Lexical Similarity	68.51	30.04	62.46	18.42	70.47	24.73	79.98	52.49	70.35	31.42
Eigenscore	66.57	25.98	73.33	34.87	79.08	31.70	81.32	56.83	75.07	37.34
HIDE	78.16	39.28	80.77	52.52	77.45	35.56	60.25	20.98	74.15	37.08
Gemma-2-9B										
Perplexity	59.69	19.95	46.14	-4.46	29.75	-4.55	40.46	-7.66	44.01	0.82
Energy	44.45	-9.16	53.56	2.39	63.18	2.08	71.93	15.07	58.28	2.60
LN-Entropy	63.10	22.58	59.07	13.84	41.64	-2.82	38.69	-7.53	50.63	6.52
Lexical Similarity	68.93	33.21	66.84	27.57	49.29	-0.60	62.08	10.82	61.79	17.75
Eigenscore	64.41	25.28	74.27	38.40	79.23	8.06	72.24	17.68	72.54	22.36
HIDE	58.11	14.86	74.57	40.93	74.12	20.87	80.94	46.90	71.94	30.89
Gemma-2-9B-Instruct										
Perplexity	59.65	12.24	52.92	3.96	30.70	-8.91	28.99	-7.91	43.07	-0.16
Energy	59.03	12.08	81.68	51.69	78.04	19.77	92.75	34.10	77.87	29.41
LN-Entropy	62.39	16.99	68.31	24.85	44.40	-3.54	36.64	-6.75	52.94	7.89
Lexical Similarity	73.28	39.48	73.55	35.35	53.91	3.30	49.55	0.13	62.57	19.57
Eigenscore	70.96	32.03	72.42	40.43	66.72	12.17	60.28	7.88	67.59	23.13
HIDE	66.08	20.79	72.62	37.79	84.00	41.94	67.92	38.57	72.65	34.77

Table C.1: AUC-ROC and PCC values using exact match as the correctness measure for faithfulness and factuality hallucination detection across four datasets. HIDE is compared with single-pass and multi-pass baselines (rows corresponding to multi-pass are highlighted in grey; 5 outputs are generated for each multi-pass method). The best result for each metric is shown in **bold**; the best among single-pass methods is underlined.

Example from SQuAD

Prompt: College sports are also popular in southern California. The UCLA Bruins and the USC Trojans both field teams in NCAA Division I in the Pac-12 Conference, and there is a longtime rivalry between the schools.

Q: The Bruins belong to which college?

A:

Ground-truth: UCLA

Llama-3-8B Output: The University of California, Los Angeles.

Input keywords: [UCLA, rivalry, southern, sports, Pac, schools, Trojans, belong, popular, Division, Q, field, longtime, in, both, The, also, is, 12, there]

Output keywords: [University, Angeles, California, Los, The, of]

Selected Input Tokens: [UCLA, rivalry, southern, sports, Pac, schools]

Selected Output Tokens: [University, Angeles, California, Los, The, of]

HIDE Score: 0.139

Example from NQ

Prompt: Answer these questions:

Q: Who played solomon in little house on the prairie?

A:

Ground-truth: Todd Bridges

Llama-3-8B Output: The actor who played Solomon in Little House on the Prairie was Jonathan Gilbert.

Input keywords: [solomon, prairie, house, questions, Q, little, who, played, A, the, Answer, in, these, on]

Output keywords: [Gilbert, Solomon, Jonathan, actor, Prairie, Little, House, who, was, played, The, the, in, on]

Selected Input Tokens: [sol, omon, pr, airie, house, questions, Q, little, who, played, A, the, Answer, in]

Selected Output Tokens: [Gilbert, Solomon, Jonathan, actor, Prairie, Little, House, who, was, played, The, the, in, on]

HIDE Score: 0.066

Example from NQ

Prompt: Answer these questions:

Q: Who plays joker in batman the dark knight?

A:

Ground-truth: Ledger

Llama-3-8B Output: Heath Ledger

Input keywords: [joker, batman, knight, questions, plays, dark, Q, A, who, Answer, these, in, the]

Output keywords: [Ledger, Heath]

Selected Input Tokens: [joker, bat]

Selected Output Tokens: [Ledger, Heath]

HIDE Score: 0.249

Example from TriviaQA

Prompt: Answer these questions:

Q: Who was the man behind The Chipmunks?

A:

Ground-truth: David Seville

Llama-3-8B Output: Ross Bagdasarian, Sr.

Input keywords: [Chipmunks, questions, Who, Q, The, the, A, behind, was, man, Answer, these]

Output keywords: [Chipmunks, Chipmunk, Bagdasarian, album, Songbook, Ross, name, first, The, the, of, Sr, A, Q, was, What]

Selected Input Tokens: [Chip, m, unks, questions, Who, Q, The, the, A, behind, was, man, Answer,

these]
Selected Output Tokens: [Chip, m, unks, Chip, m, unk, Bag, das, arian, album, Song, book, Ross, name]
HIDE Score: 0.066

Example from TriviaQA

Prompt: College sports are also popular in southern California. The UCLA Bruins and the USC Trojans both field teams in NCAA Division I in the Pac-12 Conference, and there is a longtime rivalry between the schools.
Q: What is the Japanese share index called?
A:

Ground-truth: Nikkei
Llama-3-8B Output: The Nikkei

Input keywords: [index, Japanese, share, Q, questions, A, Answer, these, the, What, called, is]
Output keywords: [Nikkei, The]
Selected Input Tokens: [index, Japanese, share, Q]
Selected Output Tokens: [Nik, ke, i, The]
HIDE Score: 0.187

Appendix E: Prevalence and Impact of Very Short Outputs

HSIC-based dependence estimation requires multiple output tokens to meaningfully assess statistical coupling between input and output representations. Consequently, HIDE is not informative when the generated output consists of a single token and is less reliable for extremely short responses.

To contextualize this limitation, Table E.1 reports the prevalence of very short model-generated outputs for Llama-3-8B across QA datasets. We observe a clear dataset-dependent pattern: TriviaQA, which emphasizes factual recall, produces a substantially higher fraction of single-token and two-token answers, whereas SQuAD and RACE more often elicit short phrases or sentence-level responses.

Although extremely short outputs represent a genuine edge case for HSIC-based dependence estimation, they do not dominate the overall evaluation and primarily affect datasets with inherently minimal answer lengths. This analysis complements the discussion in Section 8.1 and explains why the short-output failure mode does not alter the main performance trends reported in Section 6.

Dataset	% length = 1	% length ≤ 2
SQuAD	14.9	28.8
RACE	11.1	16.6
NQ	12.0	31.1
TriviaQA	36.6	71.5

Table E.1: Prevalence of very short model-generated outputs for Llama-3-8B across QA datasets.

Appendix F: Additional Results for the Attention and Norm Comparison

Across the 16,000 selected records, we retain 252 generations reaching the 256-token cap, nine empty evaluated answers, and seven no-output-state zero fallbacks (all seven in Gemma/NQ). These sets can overlap. No record is excluded from the paired metrics.

These analyses use the same fixed first-answer-line samples as Table 4, without score filtering or class balancing. Table F.1 reports ROUGE-L-based metrics; Table F.2 reports paired AUC_s differences. PCC in the detector tables uses continuous reference similarity, whereas the structural plots correlate detector scores with one another. These are distinct quantities.

Panel A: HIDE full-benchmark reference (Appendix B)

Model	Dataset	N	AUC _r	PCC _r
Llama-3-8B	SQuAD	5928	77.45	50.43
Llama-3-8B	RACE	3498	66.95	22.87
Llama-3-8B	NQ	3610	79.17	33.91
Llama-3-8B	TriviaQA	9960	66.82	35.19
Gemma-2-9B	SQuAD	5928	76.66	48.43
Gemma-2-9B	RACE	3498	57.99	16.50
Gemma-2-9B	NQ	3610	75.96	15.41
Gemma-2-9B	TriviaQA	9960	81.02	43.88

Panel B: Matched first-answer-line comparison

Model	Dataset	N_+	HIDE		Attention Ω		Norm $\tilde{\Delta}_{in}$	
			AUC _r	PCC _r	AUC _r	PCC _r	AUC _r	PCC _r
Llama-3-8B	SQuAD	922	77.32	48.97	85.57	60.71	51.24	-0.69
Llama-3-8B	RACE	875	58.08	9.35	59.88	14.16	47.00	-11.84
Llama-3-8B	NQ	280	78.58	31.79	84.00	32.58	71.44	22.36
Llama-3-8B	TriviaQA	1269	58.58	5.56	73.38	43.02	64.69	26.77
Gemma-2-9B	SQuAD	849	76.51	48.85	87.58	65.46	54.13	9.15
Gemma-2-9B	RACE	1185	58.12	16.29	50.41	-1.90	50.95	1.93
Gemma-2-9B	NQ	480	63.16	16.47	69.10	23.10	60.08	13.96
Gemma-2-9B	TriviaQA	1359	51.18	-5.81	66.30	32.04	62.86	21.04

Table F.1: Panel A reproduces HIDE values from Table B.1 of Appendix B without alteration (N : full-benchmark sample size). Panel B evaluates all three scores on the same 2,000 fixed seed-42 examples per model/dataset under first-answer-line stopping; N_+ refers only to Panel B. Because cohorts and stopping differ, Panel A is a reference, not a matched comparator for Panel B. N_+ counts answers with ROUGE-L greater than 0.5. AUC and PCC are multiplied by 100; PCC uses the continuous reference metric. Within Panel B, all scores use the same examples and labels; higher raw scores predict correctness. No direction is selected using test performance.

For actual selected count $n \geq 1$, substituting all-ones Gram matrices into the adapted score gives $c(n) = (n - 1)/n^2$ (and we define $c(0) = 0$ for the no-state fallback). Table F.3 compares this diagnostic with HIDE on exactly the same examples. The near-unit correlations and small AUC differences indicate that the configured score is dominated by selected-token count in these runs; they do not establish a nonlinear-dependence advantage. We did not retune the kernel on the test labels.

Model	Dataset	HIDE – attention	95% CI	HIDE – norm
Llama-3-8B	SQuAD	-9.01	[-11.43, -6.69]	+22.55
Llama-3-8B	RACE	-3.00	[-5.27, -0.82]	+11.69
Llama-3-8B	NQ	-5.77	[-9.66, -2.21]	+6.42
Llama-3-8B	TriviaQA	-16.67	[-19.96, -13.56]	-7.67
Gemma-2-9B	SQuAD	-10.82	[-13.37, -8.56]	+21.58
Gemma-2-9B	RACE	+6.69	[+3.70, +9.69]	+7.37
Gemma-2-9B	NQ	-5.23	[-8.69, -1.86]	+4.11
Gemma-2-9B	TriviaQA	-17.11	[-20.84, -13.51]	-11.70

Table F.2: Paired AUC_s differences in percentage points on the corrected 2,000-example cohorts. Intervals use 2,000 paired example bootstrap draws (seed 42), are pointwise, and are not adjusted for multiple comparisons.

Model	Dataset	HIDE AUC _s	Count-control AUC _s	PCC(HIDE, control)
Llama-3-8B	SQuAD	76.33	76.38	1.000000
Llama-3-8B	RACE	61.66	61.92	1.000000
Llama-3-8B	NQ	78.47	78.57	1.000000
Llama-3-8B	TriviaQA	57.01	56.58	1.000000
Gemma-2-9B	SQuAD	75.90	75.92	0.998759
Gemma-2-9B	RACE	57.93	57.83	0.998718
Gemma-2-9B	NQ	62.74	63.01	0.999911
Gemma-2-9B	TriviaQA	49.83	49.78	0.999988

Table F.3: Selected-token-count diagnostic on the same first-answer-line samples. The control is $(n_{\text{eff}} - 1)/n_{\text{eff}}^2$ for $n_{\text{eff}} \geq 1$, and zero otherwise. AUC is multiplied by 100; PCC is unscaled. This is a control for actual selected-token count, not raw output length.

Model	Dataset	Mean input tokens	Mean generated tokens
Llama-3.2-3B	SQuAD	186.47	33.44
Llama-3.2-3B	NQ	20.30	17.34
Llama-3-8B	SQuAD	186.47	21.39
Llama-3-8B	NQ	20.30	20.50
Gemma-2-9B	SQuAD	186.20	39.29
Gemma-2-9B	NQ	21.14	22.16
Gemma-2-27B	SQuAD	186.20	45.65
Gemma-2-27B	NQ	21.14	12.14

Table F.4: Mean token lengths for the timing queries in Table 6, under the benchmark stopping protocol (maximum 256 generated tokens). Means include all queries and measured repeats; repeats share deterministic token sequences. These length differences limit a parameter-only interpretation of the scaling comparison.

References

- Azaria, Amos and Tom Mitchell. 2023. The internal state of an llm knows when it’s lying. In *Findings of the Association for Computational Linguistics: EMNLP 2023*, pages 967–976.
- Bender, Emily M., Timnit Gebru, Angelina McMillan-Major, and Shmargaret Shmitchell. 2021. On the dangers of stochastic parrots: Can language models be too big? In *Proceedings of the 2021 ACM Conference on Fairness, Accountability, and Transparency*, FAccT ’21, pages 610–623, Association for Computing Machinery, New York, NY, USA.

- Brown, Tom, Benjamin Mann, Nick Ryder, Melanie Subbiah, Jared D Kaplan, Prafulla Dhariwal, Arvind Neelakantan, Pranav Shyam, Girish Sastry, Amanda Askell, Sandhini Agarwal, Ariel Herbert-Voss, Gretchen Krueger, Tom Henighan, Rewon Child, Aditya Ramesh, Daniel Ziegler, Jeffrey Wu, Clemens Winter, Chris Hesse, Mark Chen, Eric Sigler, Mateusz Litwin, Scott Gray, Benjamin Chess, Jack Clark, Christopher Berner, Sam McCandlish, Alec Radford, Ilya Sutskever, and Dario Amodei. 2020. Language models are few-shot learners. In *Advances in Neural Information Processing Systems*, volume 33, pages 1877–1901, Curran Associates, Inc.
- Chakraborty, Tanmoy and Sarah Masud. 2024. The promethean dilemma of ai at the intersection of hallucination and creativity. *Commun. ACM*, 67(10):26–28.
- Chang, Haw-Shiuan and Andrew McCallum. 2022. Softmax bottleneck makes language models unable to represent multi-mode word distributions. In *Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 8048–8073, Association for Computational Linguistics, Dublin, Ireland.
- Chen, Chao, Kai Liu, Ze Chen, Yi Gu, Yue Wu, Mingyuan Tao, Zhihang Fu, and Jieping Ye. 2024a. Inside: LLMs’ internal states retain the power of hallucination detection. In *The Twelfth International Conference on Learning Representations*.
- Chen, Jifan, Grace Kim, Aniruddh Sriram, Greg Durrett, and Eunsol Choi. 2024b. Complex claim verification with evidence retrieved in the wild. In *Proceedings of the 2024 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies (Volume 1: Long Papers)*, pages 3569–3587.
- Chern, I-Chun, Steffi Chern, Shiqi Chen, Weizhe Yuan, Kehua Feng, Chunting Zhou, Junxian He, Graham Neubig, and Pengfei Liu. 2023. Factool: Factuality detection in generative ai – a tool augmented framework for multi-task and multi-domain scenarios. *arXiv preprint arXiv:2307.13528*.
- Chiang, Cheng-Han and Hung-Yi Lee. 2023. Can large language models be an alternative to human evaluations? In *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 15607–15631.
- Chowdhery, Aakanksha, Sharan Narang, Jacob Devlin, Maarten Bosma, Gaurav Mishra, Adam Roberts, Paul Barham, Hyung Won Chung, Charles Sutton, Sebastian Gehrmann, Parker Schuh, Kensen Shi, Sasha Tsvyashchenko, Joshua Maynez, Abhishek Rao, Parker Barnes, Yi Tay, Noam Shazeer, Vinodkumar Prabhakaran, Emily Reif, Nan Du, Ben Hutchinson, Reiner Pope, James Bradbury, Jacob Austin, Michael Isard, Guy Gur-Ari, Pengcheng Yin, Toju Duke, Anselm Levskaya, Sanjay Ghemawat, Sunipa Dev, Henryk Michalewski, Xavier Garcia, Vedant Misra, Kevin Robinson, Liam Fedus, Denny Zhou, Daphne Ippolito, David Luan, Hyeontaek Lim, Barret Zoph, Alexander Spiridonov, Ryan Sepassi, David Dohan, Shivani Agrawal, Mark Omernick, Andrew M. Dai, Thanumalayan Sankaranarayanan Pillai, Marie Pellat, Aitor Lewkowycz, Erica Moreira, Rewon Child, Oleksandr Polozov, Katherine Lee, Zongwei Zhou, Xuezhi Wang, Brennan Saeta, Mark Diaz, Orhan Firat, Michele Catasta, Jason Wei, Kathy Meier-Hellstern, Douglas Eck, Jeff Dean, Slav Petrov, and Noah Fiedel. 2023. Palm: Scaling language modeling with pathways. *Journal of Machine Learning Research*, 24(240):1–113.
- Chuang, Yung-Sung, Linlu Qiu, Cheng-Yu Hsieh, Ranjay Krishna, Yoon Kim, and James Glass. 2024a. Lookback lens: Detecting and mitigating contextual hallucinations in large language models using only attention maps. In *Proceedings of EMNLP*, pages 1419–1436.
- Chuang, Yung-Sung, Yujia Xie, Hongyin Luo, Yoon Kim, James R Glass, and Pengcheng He. 2024b. Dola: Decoding by contrasting layers improves factuality in large language models. In *The Twelfth International Conference on Learning Representations*.
- Dao, Tri, Daniel Y. Fu, Stefano Ermon, Atri Rudra, and Christopher Ré. 2022. Flashattention: Fast and memory-efficient exact attention with io-awareness. *Advances in Neural Information Processing Systems*, 35.
- Duan, Hanyu, Yi Yang, and Kar Yan Tam. 2024. Do llms know about hallucination? an empirical investigation of llm’s hidden states. *arXiv preprint arXiv:2402.09733*.
- Durmus, Esin, He He, and Mona Diab. 2020. FEQA: A question answering evaluation framework for faithfulness assessment in abstractive summarization. In *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics*, pages 5055–5070, Association for Computational Linguistics, Online.
- Dziri, Nouha, Andrea Madotto, Osmar Zaiane, and Avishek Joey Bose. 2021. Neural path hunter: Reducing hallucination in dialogue systems via path grounding. In *Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing*, pages 2197–2214, Association for Computational Linguistics, Online and Punta Cana, Dominican Republic.

- Falke, Tobias, Leonardo F. R. Ribeiro, Prasetya Ajie Utama, Ido Dagan, and Iryna Gurevych. 2019. Ranking generated summaries by correctness: An interesting but challenging application for natural language inference. In *Proceedings of the 57th Annual Meeting of the Association for Computational Linguistics*, pages 2214–2220, Association for Computational Linguistics, Florence, Italy.
- Fan, Angela, Mike Lewis, and Yann Dauphin. 2018. Hierarchical neural story generation. In *Proceedings of the 56th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 889–898, Association for Computational Linguistics, Melbourne, Australia.
- Geva, Mor, Jasmijn Bastings, Katja Filippova, and Amir Globerson. 2023. Dissecting recall of factual associations in auto-regressive language models. In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*, pages 12216–12235, Association for Computational Linguistics, Singapore.
- Grattafiori, Aaron, Abhimanyu Dubey, Abhinav Jauhri, Abhinav Pandey, Abhishek Kadian, Ahmad Al-Dahle, Aiesha Letman, Akhil Mathur, Alan Schelten, Alex Vaughan, Amy Yang, Angela Fan, Anirudh Goyal, Anthony Hartshorn, Aobo Yang, Archi Mitra, Archie Sravankumar, Artem Korenev, Arthur Hinsvark, Arun Rao, Aston Zhang, Aurelien Rodriguez, Austen Gregerson, Ava Spataru, Baptiste Roziere, Bethany Biron, Binh Tang, Bobbie Chern, Charlotte Caucheteux, Chaya Nayak, Chloe Bi, Chris Marra, Chris McConnell, Christian Keller, Christophe Touret, Chunyang Wu, Corinne Wong, Cristian Canton Ferrer, Cyrus Nikolaidis, Damien Allonsius, Daniel Song, Danielle Pintz, Danny Livshits, Danny Wyatt, David Esiobu, Dhruv Choudhary, Dhruv Mahajan, Diego Garcia-Olano, Diego Perino, Dieuwke Hupkes, Egor Lakomkin, Ehab AlBadawy, Elina Lobanova, Emily Dinan, Eric Michael Smith, Filip Radenovic, Francisco Guzmán, Frank Zhang, Gabriel Synnaeve, Gabrielle Lee, Georgia Lewis Anderson, Govind Thattai, Graeme Nail, Gregoire Mialon, Guan Pang, Guillem Cucurell, Hailey Nguyen, Hannah Korevaar, Hu Xu, Hugo Touvron, Iliyan Zarov, Imanol Arrieta Ibarra, Isabel Kloumann, Ishan Misra, Ivan Evtimov, Jack Zhang, Jade Copet, Jaewon Lee, Jan Geffert, Jana Vranes, Jason Park, Jay Mahadeokar, Jeet Shah, Jelmer van der Linde, Jennifer Billock, Jenny Hong, Jenya Lee, Jeremy Fu, Jianfeng Chi, Jianyu Huang, Jiawen Liu, Jie Wang, Jiecao Yu, Joanna Bitton, Joe Spisak, Jongsoo Park, Joseph Rocca, Joshua Johnstun, Joshua Saxe, Junteng Jia, Kalyan Vasuden Alwala, Karthik Prasad, Kartikeya Upasani, Kate Plawiak, Ke Li, Kenneth Heafield, Kevin Stone, Khalid El-Arini, Krithika Iyer, Kshitiz Malik, Kuenley Chiu, Kunal Bhalla, Kushal Lakhotia, Lauren Rantala-Yeary, Laurens van der Maaten, Lawrence Chen, Liang Tan, Liz Jenkins, Louis Martin, Lovish Madaan, Lubo Malo, Lukas Blecher, Lukas Landzaat, Luke de Oliveira, Madeline Muzzi, Mahesh Pasupuleti, Mannat Singh, Manohar Paluri, Marcin Kardas, Maria Tsimpoukelli, Mathew Oldham, Mathieu Rita, Maya Pavlova, Melanie Kambadur, Mike Lewis, Min Si, Mitesh Kumar Singh, Mona Hassan, Naman Goyal, Narjes Torabi, Nikolay Bashlykov, Nikolay Bogoychev, Niladri Chatterji, Ning Zhang, Olivier Duchenne, Onur Çelebi, Patrick Alrassy, Pengchuan Zhang, Pengwei Li, Petar Vasic, Peter Weng, Prajjwal Bhargava, Pratik Dubal, Praveen Krishnan, Punit Singh Koura, Puxin Xu, Qing He, Qingxiao Dong, Ragavan Srinivasan, Raj Ganapathy, Ramon Calderer, Ricardo Silveira Cabral, Robert Stojnic, Roberta Raileanu, Rohan Maheswari, Rohit Girdhar, Rohit Patel, Romain Sauvestre, Ronnie Polidoro, Roshan Sumbaly, Ross Taylor, Ruan Silva, Rui Hou, Rui Wang, Saghar Hosseini, Sahana Chennabasappa, Sanjay Singh, Sean Bell, Seohyun Sonia Kim, Sergey Edunov, Shaoliang Nie, Sharan Narang, Sharath Rapparthi, Sheng Shen, Shengye Wan, Shruti Bhosale, Shun Zhang, Simon Vandenhende, Soumya Batra, Spencer Whitman, Sten Sootla, Stephane Collot, Suchin Gururangan, Sydney Borodinsky, Tamar Herman, Tara Fowler, Tarek Sheasha, Thomas Georgiou, Thomas Scialom, Tobias Speckbacher, Todor Mihaylov, Tong Xiao, Ujjwal Karn, Vedanuj Goswami, Vibhor Gupta, Vignesh Ramanathan, Viktor Kerkez, Vincent Gouget, Virginie Do, Vish Vogeti, Vitor Albiero, Vladan Petrovic, Weiwei Chu, Wenhan Xiong, Wenxin Fu, Whitney Meers, Xavier Martinet, Xiaodong Wang, Xiaofang Wang, Xiaoqing Ellen Tan, Xide Xia, Xinfeng Xie, Xuchao Jia, Xuewei Wang, Yaelle Goldschlag, Yashesh Gaur, Yasmine Babaei, Yi Wen, Yiwen Song, Yuchen Zhang, Yue Li, Yuning Mao, Zacharie Delpierre Coudert, Zheng Yan, Zhengxing Chen, Zoe Papakipos, Aaditya Singh, Aayushi Srivastava, Abha Jain, Adam Kelsey, Adam Shajnfeld, Adithya Gangidi, Adolfo Victoria, Ahuva Goldstand, Ajay Menon, Ajay Sharma, Alex Boesenberg, Alexei Baevski, Allie Feinstein, Amanda Kallet, Amit Sangani, Amos Teo, Anam Yunus, Andrei Lupu, Andres Alvarado, Andrew Caples, Andrew Gu, Andrew Ho, Andrew Poulton, Andrew Ryan, Ankit Ramchandani, Annie Dong, Annie Franco, Anuj Goyal, Aparajita Saraf, Arkabandhu Chowdhury, Ashley Gabriel, Ashwin Bharambe, Assaf

- Eisenman, Azadeh Yazdan, Beau James, Ben Maurer, Benjamin Leonhardi, Bernie Huang, Beth Loyd, Beto De Paola, Bhargavi Paranjape, Bing Liu, Bo Wu, Boyu Ni, Braden Hancock, Bram Wasti, Brandon Spence, Brani Stojkovic, Brian Gamido, Britt Montalvo, Carl Parker, Carly Burton, Catalina Mejia, Ce Liu, Changan Wang, Changkyu Kim, Chao Zhou, Chester Hu, Ching-Hsiang Chu, Chris Cai, Chris Tindal, Christoph Feichtenhofer, Cynthia Gao, Damon Civin, Dana Beaty, Daniel Kreymmer, Daniel Li, David Adkins, David Xu, Davide Testuggine, Delia David, Devi Parikh, Diana Liskovich, Didem Foss, Dingkan Wang, Duc Le, Dustin Holland, Edward Dowling, Eissa Jamil, Elaine Montgomery, Eleonora Presani, Emily Hahn, Emily Wood, Eric-Tuan Le, Erik Brinkman, Esteban Arcaute, Evan Dunbar, Evan Smothers, Fei Sun, Felix Kreuk, Feng Tian, Filippas Kokkinos, Firat Ozgenel, Francesco Caggioni, Frank Kanayet, Frank Seide, Gabriela Medina Florez, Gabriella Schwarz, Gada Badeer, Georgia Sweet, Gil Halpern, Grant Herman, Grigory Sizov, Guangyi, Zhang, Guna Lakshminarayanan, Hakan Inan, Hamid Shojanazeri, Han Zou, Hannah Wang, Hanwen Zha, Haroun Hameeb, Harrison Rudolph, Helen Suk, Henry Aspegren, Hunter Goldman, Hongyuan Zhan, Ibrahim Damlaj, Igor Molybog, Igor Tufanov, Ilias Leontiadis, Irina-Elena Veliche, Itai Gat, Jake Weissman, James Geboski, James Kohli, Janice Lam, Japhet Asher, Jean-Baptiste Gaya, Jeff Marcus, Jeff Tang, Jennifer Chan, Jenny Zhen, Jeremy Reizenstein, Jeremy Teboul, Jessica Zhong, Jian Jin, Jingyi Yang, Joe Cummings, Jon Carvill, Jon Shepard, Jonathan McPhie, Jonathan Torres, Josh Ginsburg, Junjie Wang, Kai Wu, Kam Hou U, Karan Saxena, Kartikay Khandelwal, Katayoun Zand, Kathy Matosich, Kaushik Veeraraghavan, Kelly Michelena, Keqian Li, Kiran Jagadeesh, Kun Huang, Kunal Chawla, Kyle Huang, Lailin Chen, Lakshya Garg, Lavender A, Leandro Silva, Lee Bell, Lei Zhang, Liangpeng Guo, Licheng Yu, Liron Moshkovich, Luca Wehrstedt, Madian Khabsa, Manav Avalani, Manish Bhatt, Martynas Mankus, Matan Hasson, Matthew Lennie, Matthias Reso, Maxim Groshev, Maxim Naumov, Maya Lathi, Meghan Keneally, Miao Liu, Michael L. Seltzer, Michal Valko, Michelle Restrepo, Mihir Patel, Mik Vyatskov, Mikayel Samvelyan, Mike Clark, Mike Macey, Mike Wang, Miquel Jubert Hermoso, Mo Metanat, Mohammad Rastegari, Munish Bansal, Nandhini Santhanam, Natascha Parks, Natasha White, Navyata Bawa, Nayan Singhal, Nick Egebo, Nicolas Usunier, Nikhil Mehta, Nikolay Pavlovich Laptev, Ning Dong, Norman Cheng, Oleg Chernoguz, Olivia Hart, Omkar Salpekar, Ozlem Kalinli, Parkin Kent, Parth Parekh, Paul Saab, Pavan Balaji, Pedro Rittner, Philip Bontrager, Pierre Roux, Piotr Dollar, Polina Zvyagina, Prashant Ratanchandani, Pritish Yuvraj, Qian Liang, Rachad Alao, Rachel Rodriguez, Rafi Ayub, Raghotham Murthy, Raghu Nayani, Rahul Mitra, Rangaprabhu Parthasarathy, Raymond Li, Rebekkah Hogan, Robin Battey, Rocky Wang, Russ Howes, Ruty Rinott, Sachin Mehta, Sachin Siby, Sai Jayesh Bondu, Samyak Datta, Sara Chugh, Sara Hunt, Sargun Dhillon, Sasha Sidorov, Satadru Pan, Saurabh Mahajan, Saurabh Verma, Seiji Yamamoto, Sharadh Ramaswamy, Shaun Lindsay, Shaun Lindsay, Sheng Feng, Shenghao Lin, Shengxin Cindy Zha, Shishir Patil, Shiva Shankar, Shuang Zhang, Shuang Zhang, Sinong Wang, Sneha Agarwal, Soji Sajuyigbe, Soumith Chintala, Stephanie Max, Stephen Chen, Steve Kehoe, Steve Satterfield, Sudarshan Govindaprasad, Sumit Gupta, Summer Deng, Sungmin Cho, Sunny Virk, Suraj Subramanian, Sy Choudhury, Sydney Goldman, Tal Remez, Tamar Glaser, Tamara Best, Thilo Koehler, Thomas Robinson, Tianhe Li, Tianjun Zhang, Tim Matthews, Timothy Chou, Tzook Shaked, Varun Vontimitta, Victoria Ajayi, Victoria Montanez, Vijai Mohan, Vinay Satish Kumar, Vishal Mangla, Vlad Ionescu, Vlad Poenaru, Vlad Tiberiu Mihailescu, Vladimir Ivanov, Wei Li, Wenchen Wang, Wenwen Jiang, Wes Bouaziz, Will Constable, Xiaocheng Tang, Xiaoqian Wu, Xiaolan Wang, Xilun Wu, Xinbo Gao, Yaniv Kleinman, Yanjun Chen, Ye Hu, Ye Jia, Ye Qi, Yenda Li, Yilin Zhang, Ying Zhang, Yossi Adi, Youngjin Nam, Yu, Wang, Yu Zhao, Yuchen Hao, Yundi Qian, Yunlu Li, Yuzi He, Zach Rait, Zachary DeVito, Zef Rosnbrick, Zhaoduo Wen, Zhenyu Yang, Zhiwei Zhao, and Zhiyu Ma. 2024. The llama 3 herd of models. *arXiv preprint arXiv:2407.21783*.
- Gretton, Arthur, Ralf Herbrich, Alexander Smola, Olivier Bousquet, and Bernhard Schölkopf. 2005. Kernel methods for measuring independence. *Journal of Machine Learning Research*, 6(70):2075–2129.
- Hernandez, Danny, Tom Brown, Tom Conerly, Nova DasSarma, Dawn Drain, Sheer El-Showk, Nelson Elhage, Zac Hatfield-Dodds, Tom Henighan, Tristan Hume, et al. 2022. Scaling laws and interpretability of learning from repeated data. *arXiv preprint arXiv:2205.10487*.
- Holtzman, Ari, Jan Buys, Li Du, Maxwell Forbes, and Yejin Choi. 2020. The curious case of neural text degeneration. In *International Conference on Learning Representations*.

- Hong, Giwon, Aryo Pradipta Gema, Rohit Saxena, Xiaotang Du, Ping Nie, Yu Zhao, Laura Perez-Beltrachini, Max Ryabinin, Xuanli He, Clémentine Fourrier, and Pasquale Minervini. 2024. The hallucinations leaderboard - an open effort to measure hallucinations in large language models. *CoRR*, abs/2404.05904.
- Huang, Lei, Weijiang Yu, Weitao Ma, Weihong Zhong, Zhangyin Feng, Haotian Wang, Qianglong Chen, Weihua Peng, Xiaocheng Feng, Bing Qin, and Ting Liu. 2025. A survey on hallucination in large language models: Principles, taxonomy, challenges, and open questions. *ACM Trans. Inf. Syst.*, 43(2).
- Huo, Siqing, Negar Arabzadeh, and Charles LA Clarke. 2023. Retrieving supporting evidence for llms generated answers. *arXiv preprint arXiv:2306.13781*.
- Ji, Ziwei, Delong Chen, Etsuko Ishii, Samuel Cahyawijaya, Yejin Bang, Bryan Wilie, and Pascale Fung. 2024. LLM internal states reveal hallucination risk faced with a query. In *Proceedings of the 7th BlackboxNLP Workshop: Analyzing and Interpreting Neural Networks for NLP*, pages 88–104, Association for Computational Linguistics, Miami, Florida, US.
- Ji, Ziwei, Nayeon Lee, Rita Frieske, Tiezheng Yu, Dan Su, Yan Xu, Etsuko Ishii, Ye Jin Bang, Andrea Madotto, and Pascale Fung. 2023. Survey of hallucination in natural language generation. *ACM Comput. Surv.*, 55(12).
- Joshi, Mandar, Eunsol Choi, Daniel S Weld, and Luke Zettlemoyer. 2017. Triviaqa: A large scale distantly supervised challenge dataset for reading comprehension. In *Proceedings of the 55th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 1601–1611.
- Kadavath, Saurav, Tom Conerly, Amanda Askell, Tom Henighan, Dawn Drain, Ethan Perez, Nicholas Schiefer, Zac Hatfield-Dodds, Nova DasSarma, Eli Tran-Johnson, Scott Johnston, Sheer El-Showk, Andy Jones, Nelson Elhage, Tristan Hume, Anna Chen, Yuntao Bai, Sam Bowman, Stanislav Fort, Deep Ganguli, Danny Hernandez, Josh Jacobson, Jackson Kernion, Shauna Kravec, Liane Lovitt, Kamal Ndousse, Catherine Olsson, Sam Ringer, Dario Amodei, Tom Brown, Jack Clark, Nicholas Joseph, Ben Mann, Sam McCandlish, Chris Olah, and Jared Kaplan. 2022. Language models (mostly) know what they know. *arXiv preprint arXiv:2207.05221*.
- Kasai, Jungo, Keisuke Sakaguchi, Ronan Le Bras, Akari Asai, Xinyan Yu, Dragomir Radev, Noah A Smith, Yejin Choi, Kentaro Inui, et al. 2023. Realtime qa: What’s the answer right now? *Advances in neural information processing systems*, 36:49025–49043.
- Kobayashi, Goro, Tatsuki Kuribayashi, Sho Yokoi, and Kentaro Inui. 2020. Attention is not only a weight: Analyzing transformers with vector norms. In *Proceedings of EMNLP*, pages 7057–7075.
- Kwiatkowski, Tom, Jennimaria Palomaki, Olivia Redfield, Michael Collins, Ankur Parikh, Chris Alberti, Danielle Epstein, Illia Polosukhin, Jacob Devlin, Kenton Lee, Kristina Toutanova, Llion Jones, Matthew Kelcey, Ming-Wei Chang, Andrew M. Dai, Jakob Uszkoreit, Quoc Le, and Slav Petrov. 2019. Natural questions: A benchmark for question answering research. *Transactions of the Association for Computational Linguistics*, 7:452–466.
- Laban, Philippe, Wojciech Kryściński, Divyansh Agarwal, Alexander R Fabbri, Caiming Xiong, Shafiq Joty, and Chien-Sheng Wu. 2023. Llms as factual reasoners: Insights from existing benchmarks and beyond. *arXiv preprint arXiv:2305.14540*.
- Ladhak, Faisal, Esin Durmus, Mirac Suzgun, Tianyi Zhang, Dan Jurafsky, Kathleen McKeown, and Tatsunori Hashimoto. 2023. When do pre-training biases propagate to downstream tasks? a case study in text summarization. In *Proceedings of the 17th Conference of the European Chapter of the Association for Computational Linguistics*, pages 3206–3219, Association for Computational Linguistics, Dubrovnik, Croatia.
- Lai, Guokun, Qizhe Xie, Hanxiao Liu, Yiming Yang, and Eduard Hovy. 2017. RACE: Large-scale ReAding comprehension dataset from examinations. In *Proceedings of the 2017 Conference on Empirical Methods in Natural Language Processing*, pages 785–794, Association for Computational Linguistics, Copenhagen, Denmark.
- Lee, Katherine, Daphne Ippolito, Andrew Nystrom, Chiyuan Zhang, Douglas Eck, Chris Callison-Burch, and Nicholas Carlini. 2022. Deduplicating training data makes language models better. In *Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 8424–8445, Association for Computational Linguistics, Dublin, Ireland.
- Lee, Minhyeok. 2023. A mathematical investigation of hallucination and creativity in gpt models. *Mathematics*, 11(10).

- Li, Zuchao, Shitou Zhang, Hai Zhao, Yifei Yang, and Dongjie Yang. 2023. Batgpt: A bidirectional autoregressive talker from generative pre-trained transformer. *arXiv preprint arXiv:2307.00360*.
- Liang, Percy, Rishi Bommasani, Tony Lee, Dimitris Tsipras, Dilara Soylu, Michihiro Yasunaga, Yian Zhang, Deepak Narayanan, Yuhuai Wu, Ananya Kumar, Benjamin Newman, Binhang Yuan, Bobby Yan, Ce Zhang, Christian Cosgrove, Christopher D. Manning, Christopher Ré, Diana Acosta-Navas, Drew A. Hudson, Eric Zelikman, Esin Durmus, Faisal Ladhak, Frieda Rong, Hongyu Ren, Huaxiu Yao, Jue Wang, Keshav Santhanam, Laurel Orr, Lucia Zheng, Mert Yuksekgonul, Mirac Suzgun, Nathan Kim, Neel Guha, Niladri Chatterji, Omar Khattab, Peter Henderson, Qian Huang, Ryan Chi, Sang Michael Xie, Shibani Santurkar, Surya Ganguli, Tatsunori Hashimoto, Thomas Icard, Tianyi Zhang, Vishrav Chaudhary, William Wang, Xuechen Li, Yifan Mai, Yuhui Zhang, and Yuta Koreeda. 2023. Holistic evaluation of language models. *Transactions on Machine Learning Research*.
- Lin, Chin-Yew. 2004. ROUGE: A package for automatic evaluation of summaries. In *Text Summarization Branches Out*, pages 74–81, Association for Computational Linguistics, Barcelona, Spain.
- Lin, Stephanie, Jacob Hilton, and Owain Evans. 2022. TruthfulQA: Measuring how models mimic human falsehoods. In *Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 3214–3252, Association for Computational Linguistics, Dublin, Ireland.
- Lin, Zhen, Shubhendu Trivedi, and Jimeng Sun. 2024. Generating with confidence: Uncertainty quantification for black-box large language models. *Transactions on Machine Learning Research*.
- Lin, Zi, Jeremiah Zhe Liu, and Jingbo Shang. 2022. Towards collaborative neural-symbolic graph semantic parsing via uncertainty. In *Findings of the Association for Computational Linguistics: ACL 2022*, pages 4160–4173, Association for Computational Linguistics, Dublin, Ireland.
- Liu, Bingbin, Jordan Ash, Surbhi Goel, Akshay Krishnamurthy, and Cyril Zhang. 2023. Exposing attention glitches with flip-flop language modeling. *Advances in Neural Information Processing Systems*, 36:25549–25583.
- Liu, Weitang, Xiaoyun Wang, John Owens, and Yixuan Li. 2020. Energy-based Out-of-distribution Detection. In *Advances in Neural Information Processing Systems*, volume 33, pages 21464–21475, Curran Associates, Inc.
- Liu, Yijin, Xianfeng Zeng, Chenze Shao, Fandong Meng, and Jie Zhou. 2024. Instruction position matters in sequence generation with large language models. In *Findings of the Association for Computational Linguistics ACL 2024*, pages 11652–11663.
- Luo, Zheheng, Qianqian Xie, and Sophia Ananiadou. 2023. Chatgpt as a factual inconsistency evaluator for text summarization. *arXiv preprint arXiv:2303.15621*.
- Maleki, Negar, Balaji Padmanabhan, and Kaushik Dutta. 2024. AI Hallucinations: A Misnomer Worth Clarifying. In *2024 IEEE Conference on Artificial Intelligence (CAI)*, pages 133–138, IEEE Computer Society, Los Alamitos, CA, USA.
- Malinin, Andrey and Mark Gales. 2021. Uncertainty estimation in autoregressive structured prediction. In *International Conference on Learning Representations*.
- Manakul, Potsawee, Adian Liustie, and Mark Gales. 2023. Selfcheckgpt: Zero-resource black-box hallucination detection for generative large language models. In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*, pages 9004–9017.
- Maynez, Joshua, Shashi Narayan, Bernd Bohnet, and Ryan McDonald. 2020. On faithfulness and factuality in abstractive summarization. In *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics*, pages 1906–1919, Association for Computational Linguistics, Online.
- Miao, Ning, Yee Whye Teh, and Tom Rainforth. 2023. Selfcheck: Using llms to zero-shot check their own step-by-step reasoning.
- Min, Sewon, Kalpesh Krishna, Xixi Lyu, Mike Lewis, Wen-tau Yih, Pang Koh, Mohit Iyyer, Luke Zettlemoyer, and Hannaneh Hajishirzi. 2023. Factscore: Fine-grained atomic evaluation of factual precision in long form text generation. In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*, pages 12076–12100.
- Nan, Feng, Ramesh Nallapati, Zhiguo Wang, Cicero Nogueira dos Santos, Henghui Zhu, Dejjiao Zhang, Kathleen McKeown, and Bing Xiang. 2021. Entity-level factual consistency of abstractive text summarization. In *Proceedings of the 16th Conference of the European Chapter of the Association for Computational Linguistics: Main Volume*, pages 2727–2733, Association for Computational Linguistics, Online.

- Narayanan Venkit, Pranav, Sanjana Gautam, Ruchi Panchanadikar, Ting-Hao Huang, and Shomir Wilson. 2023. Nationality bias in text generation. In *Proceedings of the 17th Conference of the European Chapter of the Association for Computational Linguistics*, pages 116–122, Association for Computational Linguistics, Dubrovnik, Croatia.
- Obeso, Oscar, Andy Arditi, Javier Ferrando, Joshua Freeman, Cameron Holmes, and Neel Nanda. 2025. Real-time detection of hallucinated entities in long-form generation.
- Onoe, Yasumasa, Michael Zhang, Eunsol Choi, and Greg Durrett. 2022. Entity cloze by date: what lms know about unseen entities. In *Findings of the Association for Computational Linguistics: NAACL 2022*, pages 693–702, Association for Computational Linguistics, Seattle, United States.
- Paullada, Amandalynne, Inioluwa Deborah Raji, Emily M Bender, Emily Denton, and Alex Hanna. 2021. Data and its (dis) contents: A survey of dataset development and use in machine learning research. *Patterns*, 2(11).
- Rajpurkar, Pranav, Jian Zhang, Konstantin Lopyrev, and Percy Liang. 2016. Squad: 100,000+ questions for machine comprehension of text. In *Proceedings of the 2016 Conference on Empirical Methods in Natural Language Processing*, pages 2383–2392.
- Reimers, Nils and Iryna Gurevych. 2019. Sentence-BERT: Sentence embeddings using Siamese BERT-networks. In *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing (EMNLP-IJCNLP)*, pages 3982–3992, Association for Computational Linguistics, Hong Kong, China.
- Ren, Jie, Jiaming Luo, Yao Zhao, Kundan Krishna, Mohammad Saleh, Balaji Lakshminarayanan, and Peter J Liu. 2023. Out-of-distribution detection and selective generation for conditional language models. In *The Eleventh International Conference on Learning Representations*.
- Scialom, Thomas, Paul-Alexis Dray, Sylvain Lamprier, Benjamin Piwowarski, Jacopo Staiano, Alex Wang, and Patrick Gallinari. 2021. QuestEval: Summarization asks for fact-based evaluation. In *Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing*, pages 6594–6604, Association for Computational Linguistics, Online and Punta Cana, Dominican Republic.
- Sharma, Mrinank, Meg Tong, Tomasz Korbak, David Duvenaud, Amanda Askell, Samuel R Bowman, Newton Cheng, Esin Durmus, Zac Hatfield-Dodds, Scott R Johnston, et al. 2024. Towards understanding sycophancy in language models. In *12th International Conference on Learning Representations, ICLR 2024*.
- Shi, Weijia, Xiaochuang Han, Mike Lewis, Yulia Tsvetkov, Luke Zettlemoyer, and Wen-tau Yih. 2024. Trusting your evidence: Hallucinate less with context-aware decoding. In *Proceedings of the 2024 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies (Volume 2: Short Papers)*, pages 783–791.
- Simhi, Adi, Itay Itzhak, Fazl Barez, Gabriel Stanovsky, and Yonatan Belinkov. 2025. Trust me, i'm wrong: High-certainty hallucinations in llms. *arXiv preprint arXiv:2502.12964*.
- Skean, Oscar, Md Rifat Arefin, Dan Zhao, Niket Nikul Patel, Jalal Naghiyev, Yann LeCun, and Ravid Shwartz-Ziv. 2025. Layer by layer: Uncovering hidden representations in language models. In *Forty-second International Conference on Machine Learning*.
- Song, Le, Alex Smola, Arthur Gretton, Justin Bedo, and Karsten Borgwardt. 2012. Feature selection via dependence maximization. *Journal of Machine Learning Research*, 13(47):1393–1434.
- Sriperumbudur, Bharath K., Arthur Gretton, Kenji Fukumizu, Bernhard Schölkopf, and Gert R.G. Lanckriet. 2010. Hilbert space embeddings and metrics on probability measures. *Journal of Machine Learning Research*, 11(50):1517–1561.
- Sriperumbudur, BK., A. Gretton, K. Fukumizu, G. Lanckriet, and B. Schölkopf. 2008. Injective hilbert space embeddings of probability measures. In *Proceedings of the 21st Annual Conference on Learning Theory*, pages 111–122, Max-Planck-Gesellschaft, Omnipress, Madison, WI, USA.
- Team, Gemma, Morgane Riviere, Shreya Pathak, Pier Giuseppe Sessa, Cassidy Hardin, Surya Bhupatiraju, Léonard Hussenot, Thomas Mesnard, Bobak Shahriari, Alexandre Ramé, et al. 2024. Gemma 2: Improving open language models at a practical size. *arXiv preprint arXiv:2408.00118*.
- Touvron, Hugo, Louis Martin, Kevin Stone, Peter Albert, Amjad Almahairi, Yasmine Babaei, Nikolay Bashlykov, Soumya Batra, Prajwal Bhargava, Shruti Bhosale, Dan Bikel, Lukas Blecher, Cristian Canton Ferrer, Moya Chen, Guillem Cucurull, David Esiobu, Jude Fernandes, Jeremy Fu, Wenyin Fu, Brian Fuller, Cynthia Gao, Vedanuj Goswami, Naman Goyal, Anthony Hartshorn, Saghar Hosseini, Rui Hou, Hakan Inan, Marcin Kardas, Viktor Kerkez, Madian

- Khabsa, Isabel Kloumann, Artem Korenev, Punit Singh Koura, Marie-Anne Lachaux, Thibaut Lavril, Jenya Lee, Diana Liskovich, Yinghai Lu, Yuning Mao, Xavier Martinet, Todor Mihaylov, Pushkar Mishra, Igor Molybog, Yixin Nie, Andrew Poulton, Jeremy Reizenstein, Rashi Rungta, Kalyan Saladi, Alan Schelten, Ruan Silva, Eric Michael Smith, Ranjan Subramanian, Xiaoqing Ellen Tan, Binh Tang, Ross Taylor, Adina Williams, Jian Xiang Kuan, Puxin Xu, Zheng Yan, Iliyan Zarov, Yuchen Zhang, Angela Fan, Melanie Kambadur, Sharan Narang, Aurelien Rodriguez, Robert Stojnic, Sergey Edunov, and Thomas Scialom. 2023. Llama 2: Open foundation and fine-tuned chat models. *arXiv preprint arXiv:2307.09288*.
- Wang, Chaojun and Rico Sennrich. 2020. On exposure bias, hallucination and domain shift in neural machine translation. In *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics*, pages 3544–3552, Association for Computational Linguistics, Online.
- Weidinger, Laura, Jonathan Uesato, Maribeth Rauh, Conor Griffin, Po-Sen Huang, John Mellor, Amelia Glaese, Myra Cheng, Borja Balle, Atoosa Kasirzadeh, Courtney Biles, Sasha Brown, Zac Kenton, Will Hawkins, Tom Stepleton, Abeba Birhane, Lisa Anne Hendricks, Laura Rimell, William Isaac, Julia Haas, Sean Legassick, Geoffrey Irving, and Iason Gabriel. 2022. Taxonomy of risks posed by language models. In *Proceedings of the 2022 ACM Conference on Fairness, Accountability, and Transparency, FAccT '22*, pages 214–229, Association for Computing Machinery, New York, NY, USA.
- Xiong, Miao, Zhiyuan Hu, Xinyang Lu, YIFEI LI, Jie Fu, Junxian He, and Bryan Hooi. 2024. Can llms express their uncertainty? an empirical evaluation of confidence elicitation in llms. In *The Twelfth International Conference on Learning Representations*.
- Yang, Zhilin, Zihang Dai, Ruslan Salakhutdinov, and William W Cohen. 2018. Breaking the softmax bottleneck: A high-rank rnn language model. In *International Conference on Learning Representations*.
- Zhang, Yue, Yafu Li, Leyang Cui, Deng Cai, Lema Liu, Tingchen Fu, Xinting Huang, Enbo Zhao, Yu Zhang, Yulong Chen, Longyue Wang, Anh Tuan Luu, Wei Bi, Freda Shi, and Shuming Shi. 2023. Siren’s song in the ai ocean: A survey on hallucination in large language models. *arXiv preprint arXiv:2309.01219*.
- Zheng, Shen, Jie Huang, and Kevin Chen-Chuan Chang. 2023. Why does chatgpt fall short in providing truthful answers? *arXiv preprint arXiv:2304.10513*.